

AN INTRODUCTION TO DEEP LEARNING FOR ASTRONOMY

Marc Huertas-Company

Jerusalem, March 2nd, 3rd 2019



institut
universitaire
de France



REFERENCES

SEVERAL SLIDES / INFOS SHOWN HERE ARE INSPIRED/
TAKEN FROM OTHER WORKS / COURSES FOUND ONLINE

- Deep Learning: Do-It-Yourself! [Bursuc, Krzakala, Lelarge]
- DEEPMLEARNING.AI [COURSERA, Ng, Bensouda, Katanforoosh]
- MACHINE LEARNING LECTURES [Keck]
- EPFL DEEP LEARNING COURSE [Fleuret]

Thanks to all of them!

SOME PRELIMINARY NOTES

I AM NOT A MACHINE LEARNING RESEARCHER

SOME PRELIMINARY NOTES

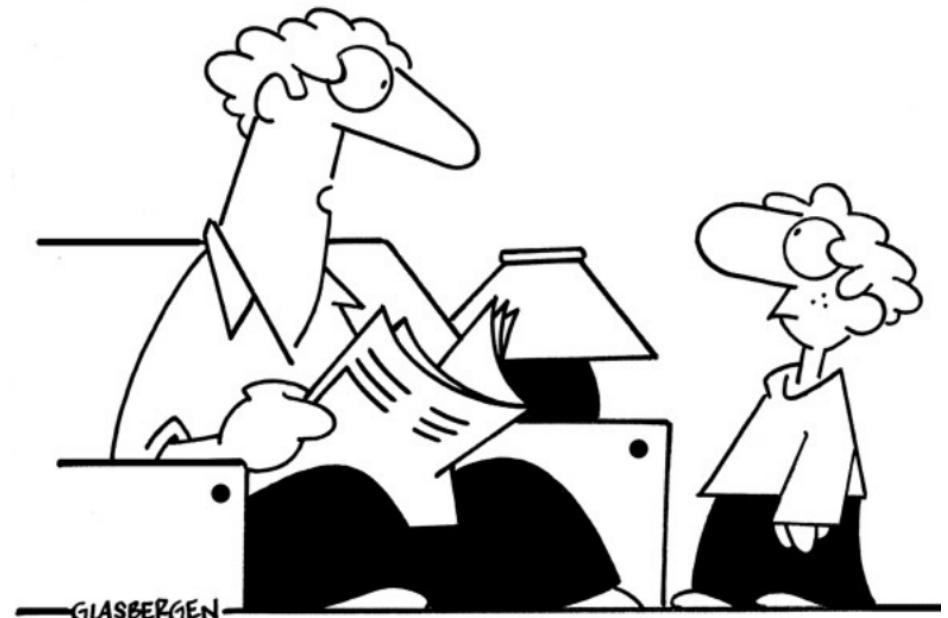
I AM NOT A MACHINE LEARNING RESEARCHER

ONLY AN ASTRONOMER WHO HAS BEEN USING MACHINE
LEARNING FOR THE LAST ~14 YEARS FOR MY RESEARCH

THESE LECTURES ARE INTENDED TO PROVIDE A **GLOBAL**
UNDERSTANDING OF HOW AI TECHNIQUES WORK AND
ESPECIALLY **HOW TO USE THEM FOR YOUR RESEARCH**

WHAT ARE WE GOING TO LEARN?

data-science
pattern-recognition
artificial-intelligence
database
data
big-data machine
data-mining
learning
clustering



*"Artificial intelligence is when you get a college degree,
but you're still stupid when you graduate."*

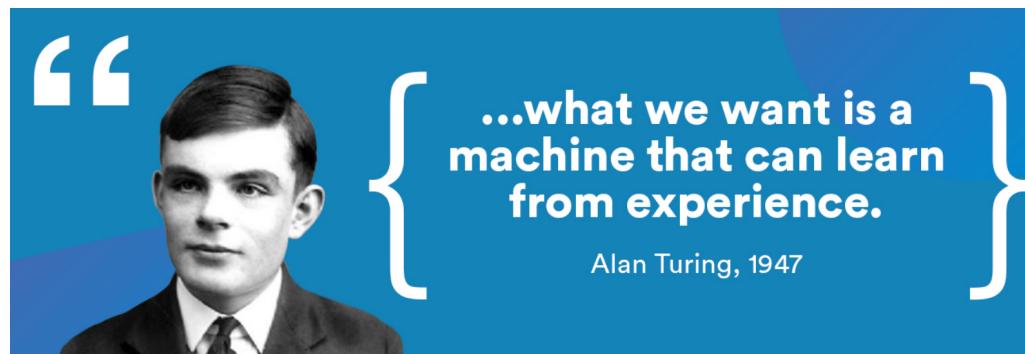
WHAT ARE WE GOING TO LEARN?

data-science
pattern-recognition
artificial-intelligence
database
da
big-data machine learning
data-mining
learning
clustering

A BUNCH OF
SOMETIMES
CONFUSING
TERMS...

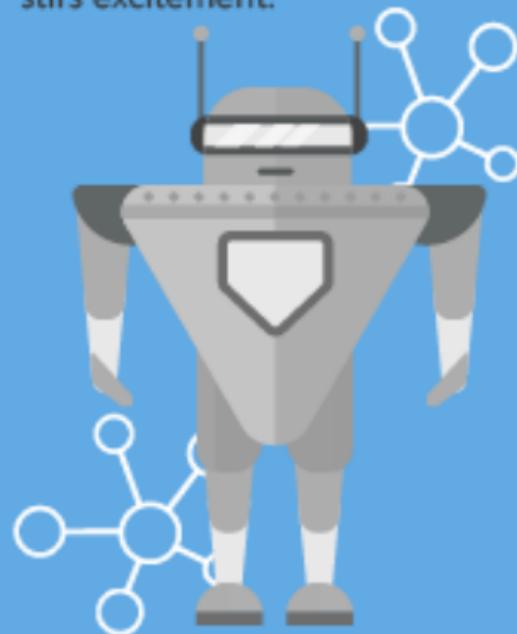


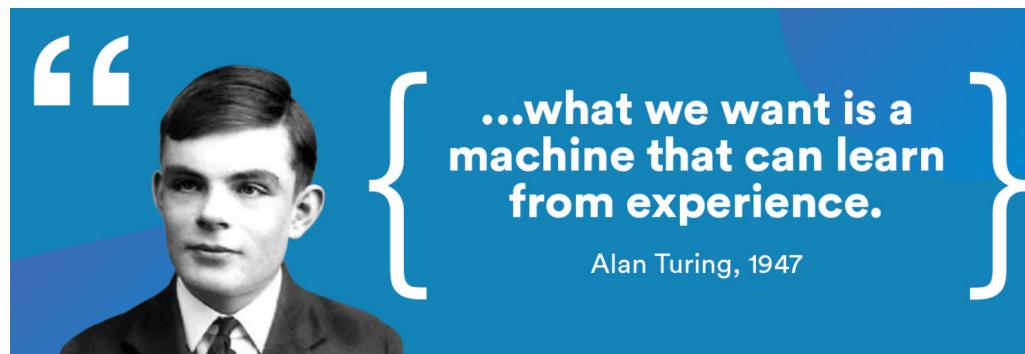
*"Artificial intelligence is when you get a college degree,
but you're still stupid when you graduate."*



ARTIFICIAL INTELLIGENCE

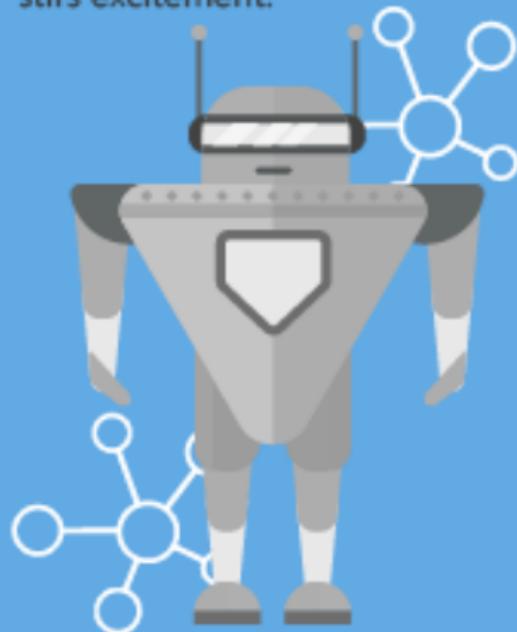
Early artificial intelligence stirs excitement.





ARTIFICIAL INTELLIGENCE

Early artificial intelligence stirs excitement.



1950's

1960's

1970's

1980's

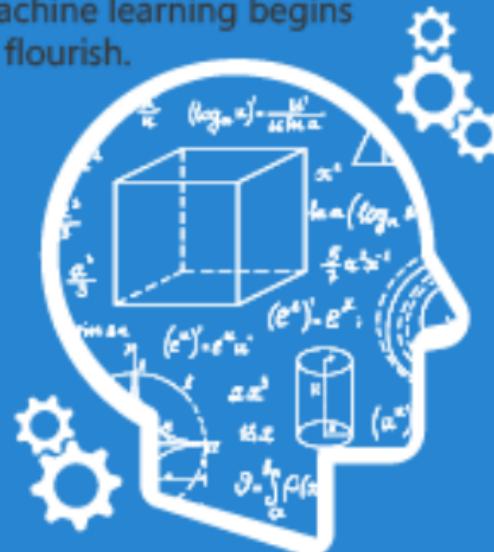
1990's

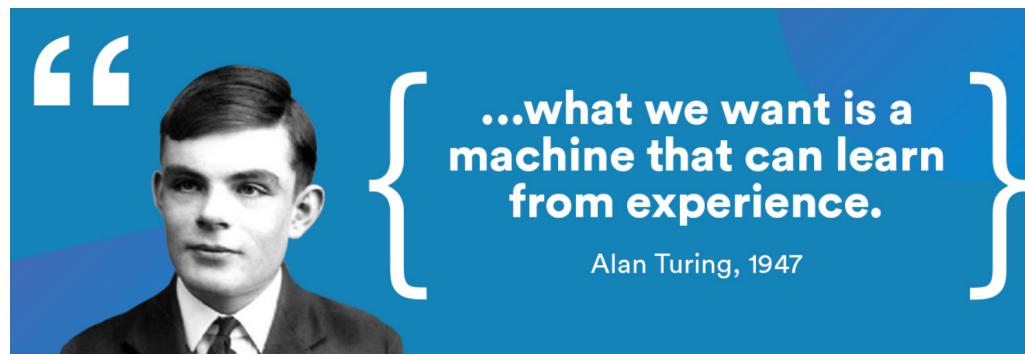
2000's

2010's

MACHINE LEARNING

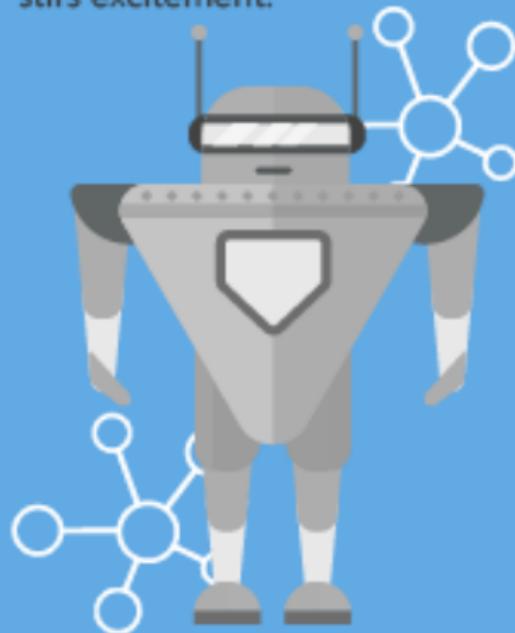
Machine learning begins to flourish.





ARTIFICIAL INTELLIGENCE

Early artificial intelligence stirs excitement.



1950's

1960's

1970's

1980's

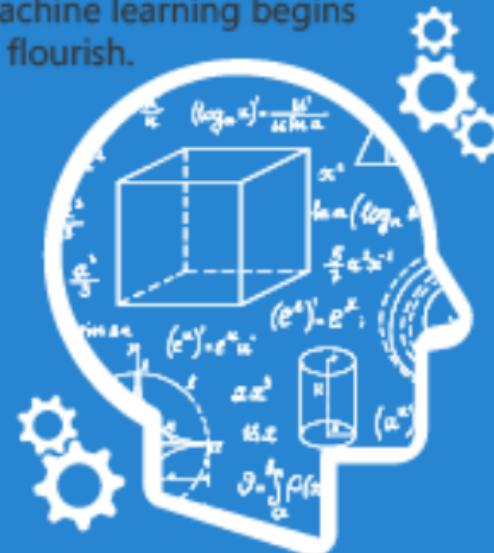
1990's

2000's

2010's

MACHINE LEARNING

Machine learning begins to flourish.



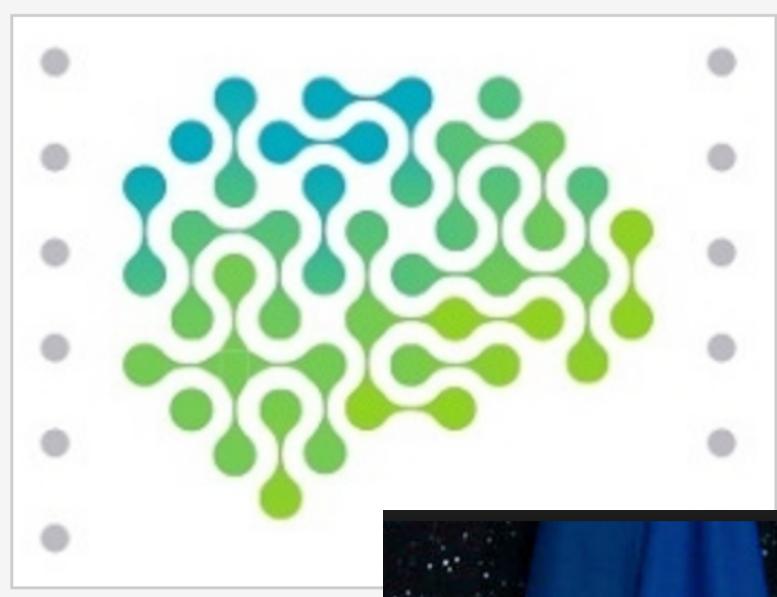
DEEP LEARNING

Deep learning breakthroughs drive AI boom.



AI FEVER?

AN AMAZING MEDIA ATTENTION

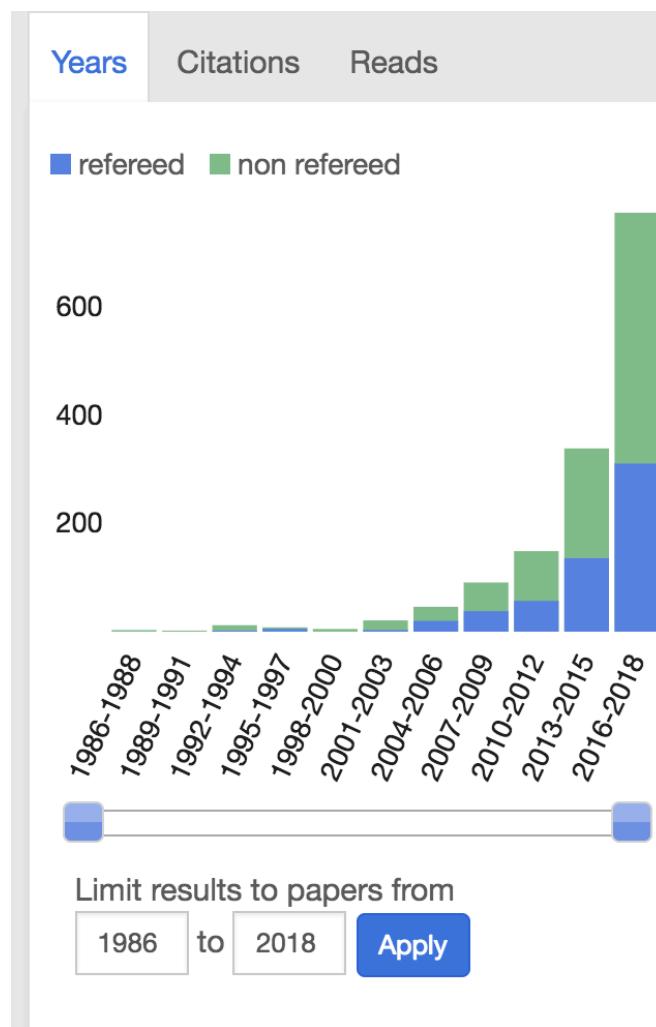


Le CNRS, Inria, l'université PSL et les entreprises Amazon, Criteo, Facebook, Faurecia, Google, Microsoft, NAVER LABS, Nokia Bell Labs, le Groupe PSA, SUEZ et Valeo font converger intérêts académiques et industriels et s'unissent pour créer, à Paris, l'Institut PRAIRIE dont l'objectif est de devenir une référence internationale de l'intelligence artificielle.

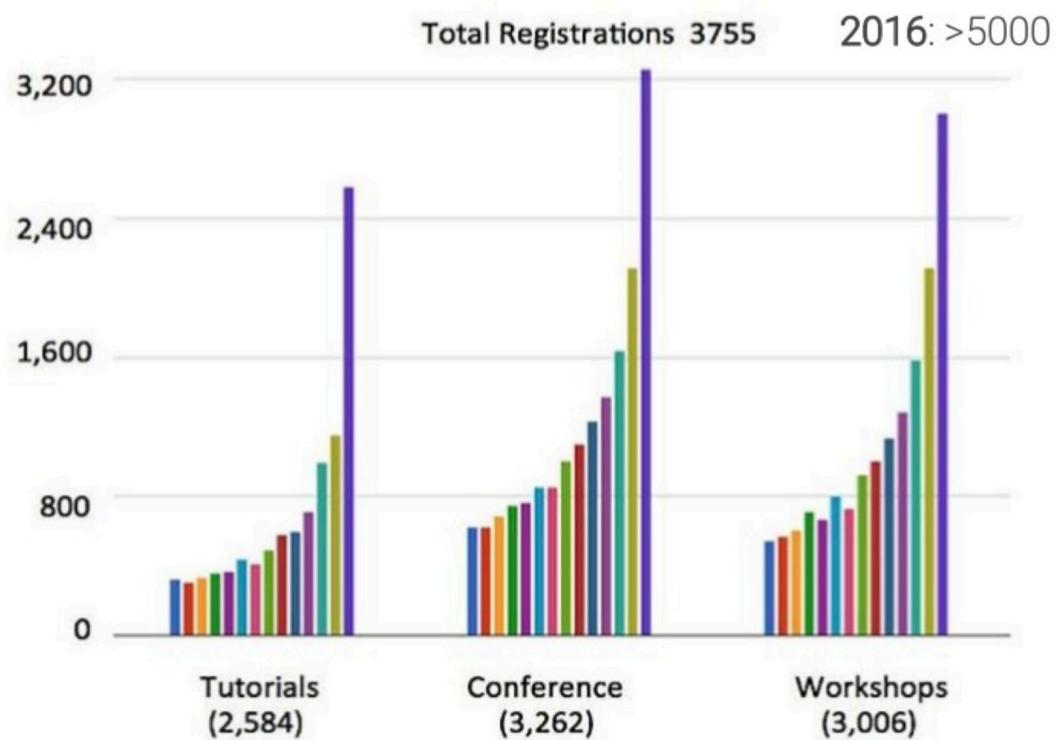


AI FEVER?

PUBLICATIONS (ADS)

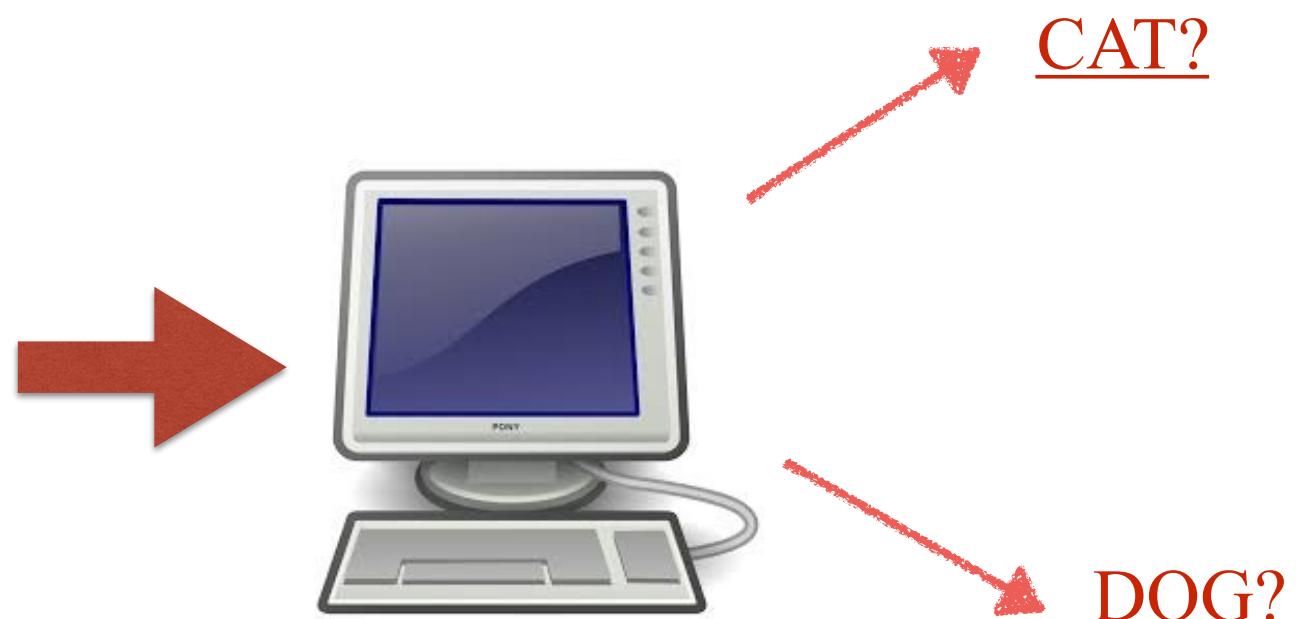


CONFERENCES



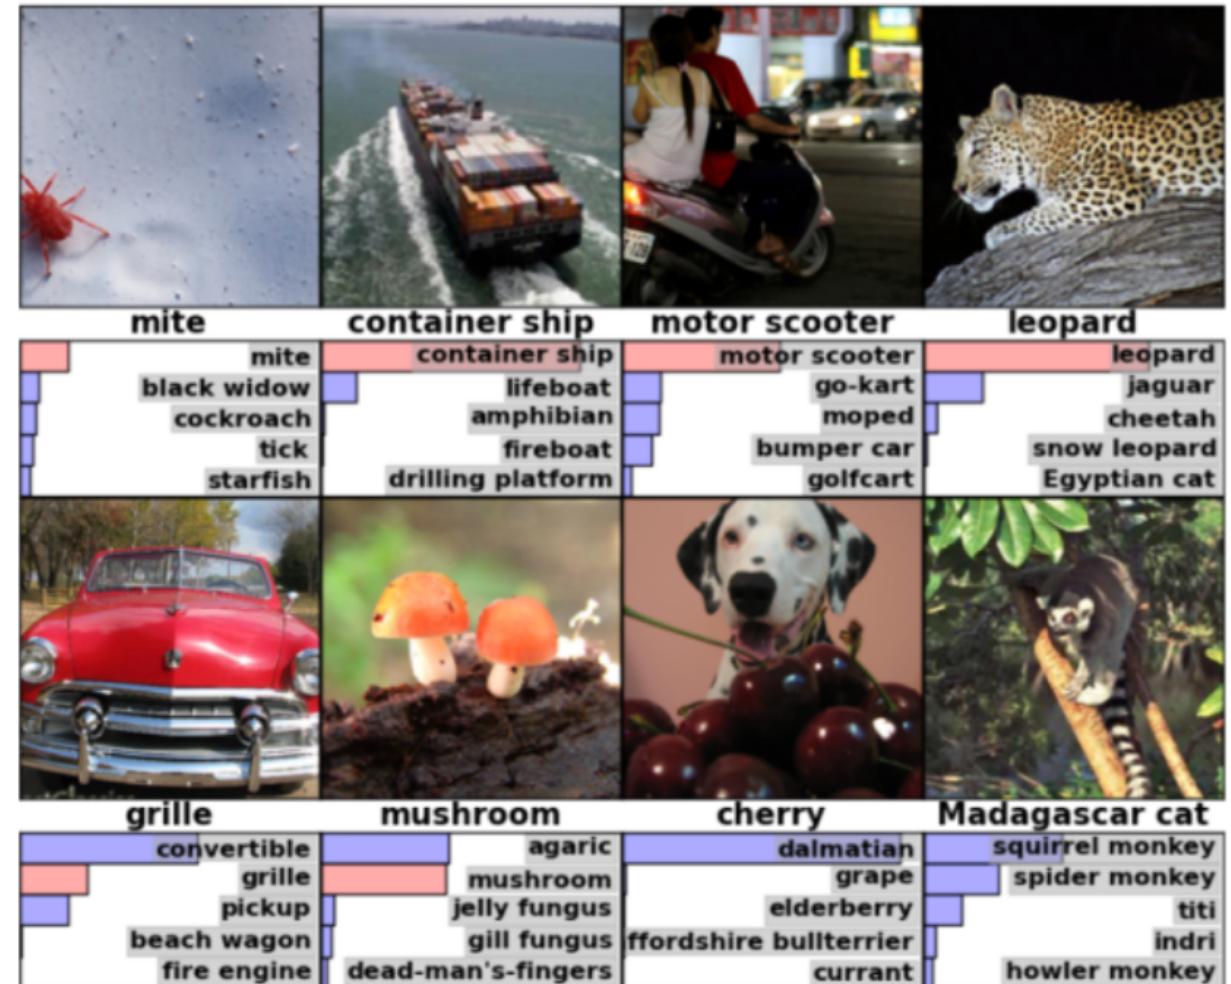
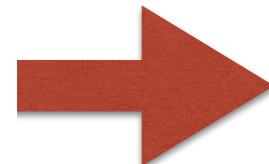
Source

BEFORE 2012....



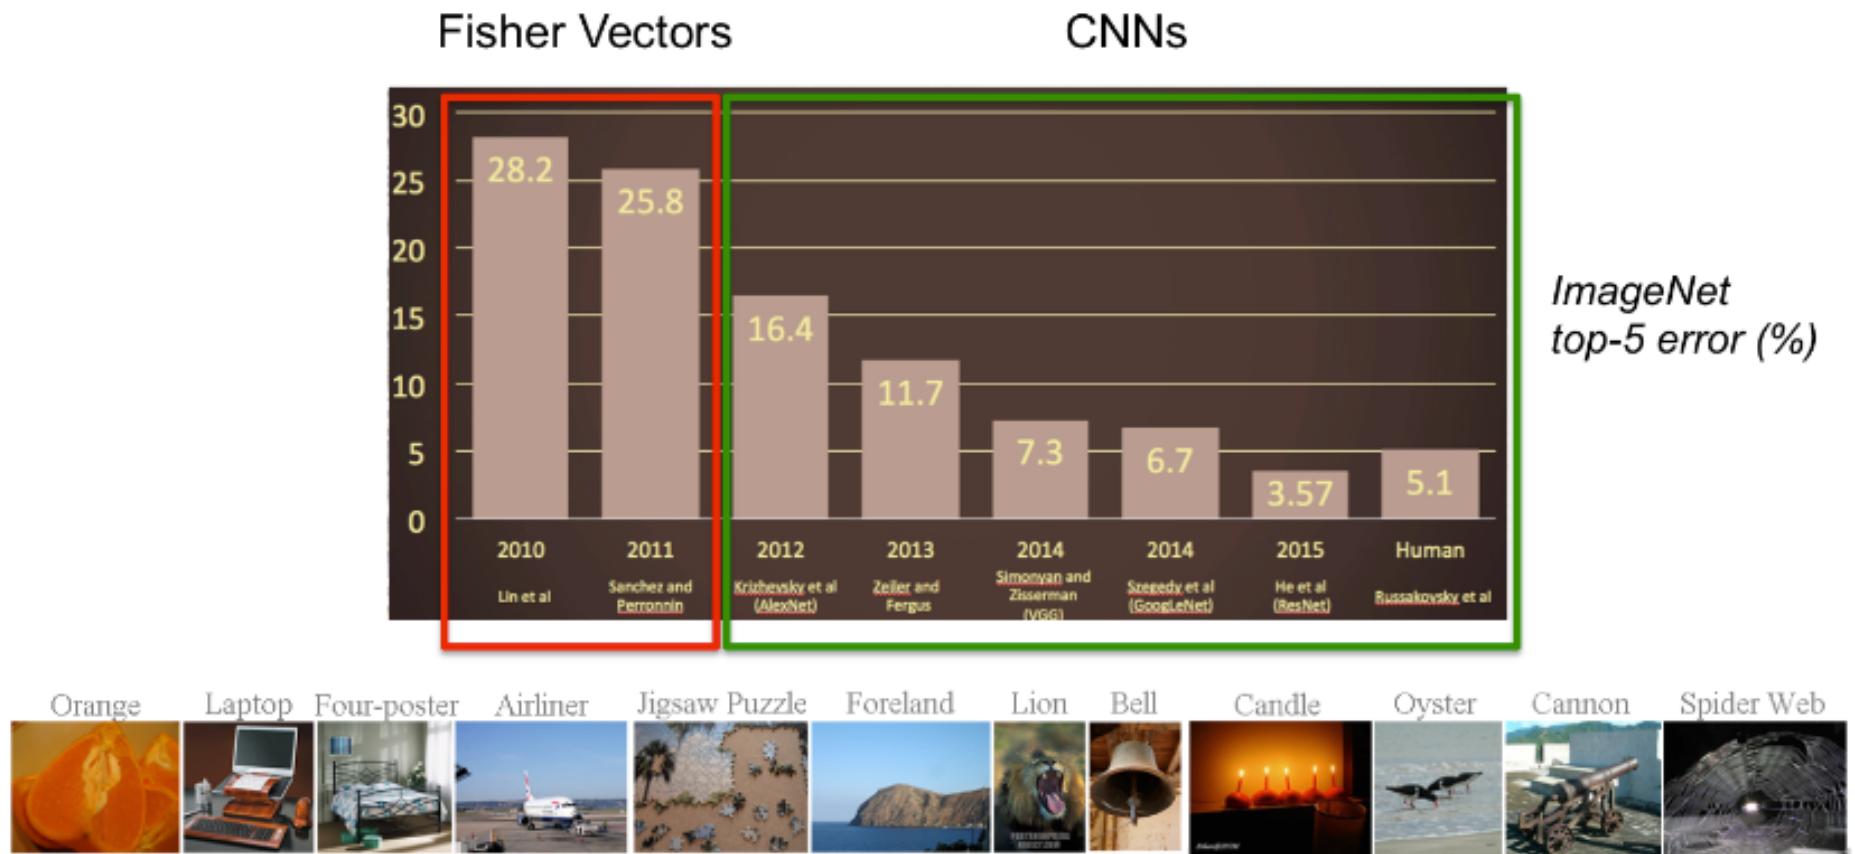
TRIVIAL HUMAN TASKS REMAINED
CHALLENGING FOR COMPUTERS

AFTER 2012



IT HAS BECOME TRIVIAL....

THIS IS A CHANGE OF PARADIGM!



ONE OF THE MAIN REASONS OF THIS
BREAKTHROUGH IS THE AVAILABILITY OF VERY
LARGE DATASETS TO LEARN



COMBINED WITH THE TECHNOLOGY TO
PROCESS ALL THIS DATA



ONE OF THE MAIN REASONS OF THIS
BREAKTHROUGH IS THE AVAILABILITY OF VERY
LARGE DATASETS TO LEARN

HOWEVER THERE HAS NOT BEEN A MAJOR
REVOLUTIONARY IDEA



WHAT ARE WE GOING TO LEARN?

BASICS OF CLASSICAL MACHINE LEARNING

BASICS OF DEEP LEARNING
(BOTH SUPERVISED AND UNSUPERVISED)

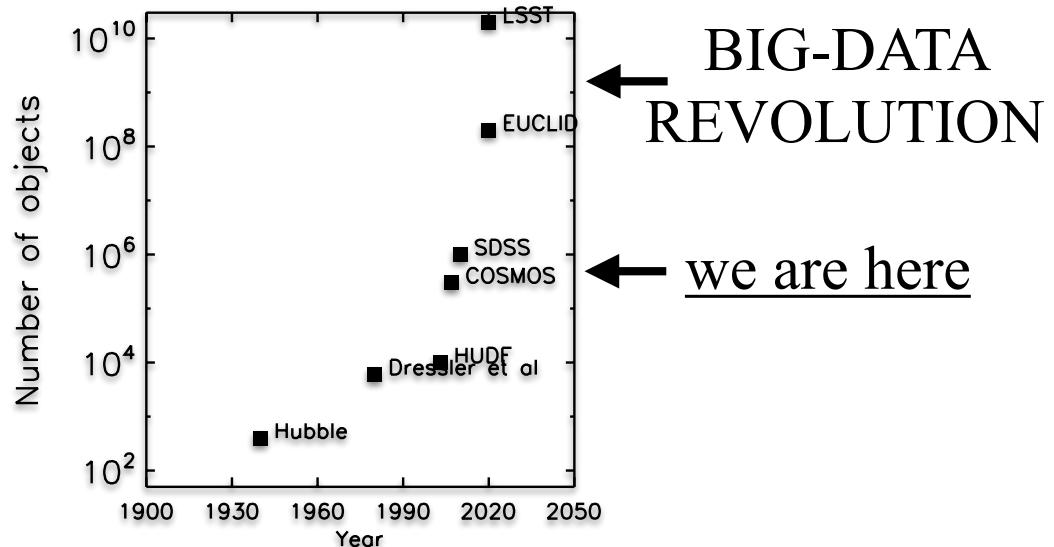
HOPING THAT THIS WOULD BE USEFUL FOR YOUR
RESEARCH!

(Apologies in advance for biases on Extra-Galactic Science +
imaging)

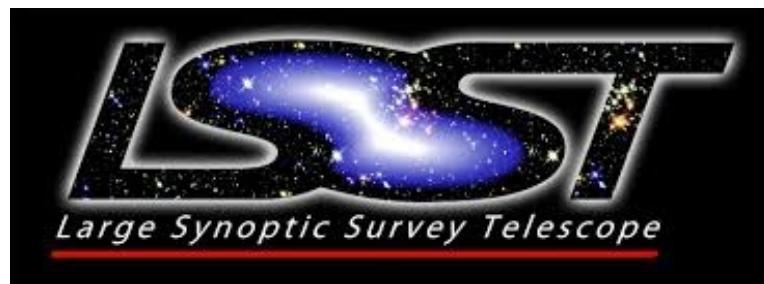
WHY DO WE NEED THESE TOOLS IN ASTRONOMY?

WHY DO WE NEED THESE TOOLS IN ASTRONOMY?

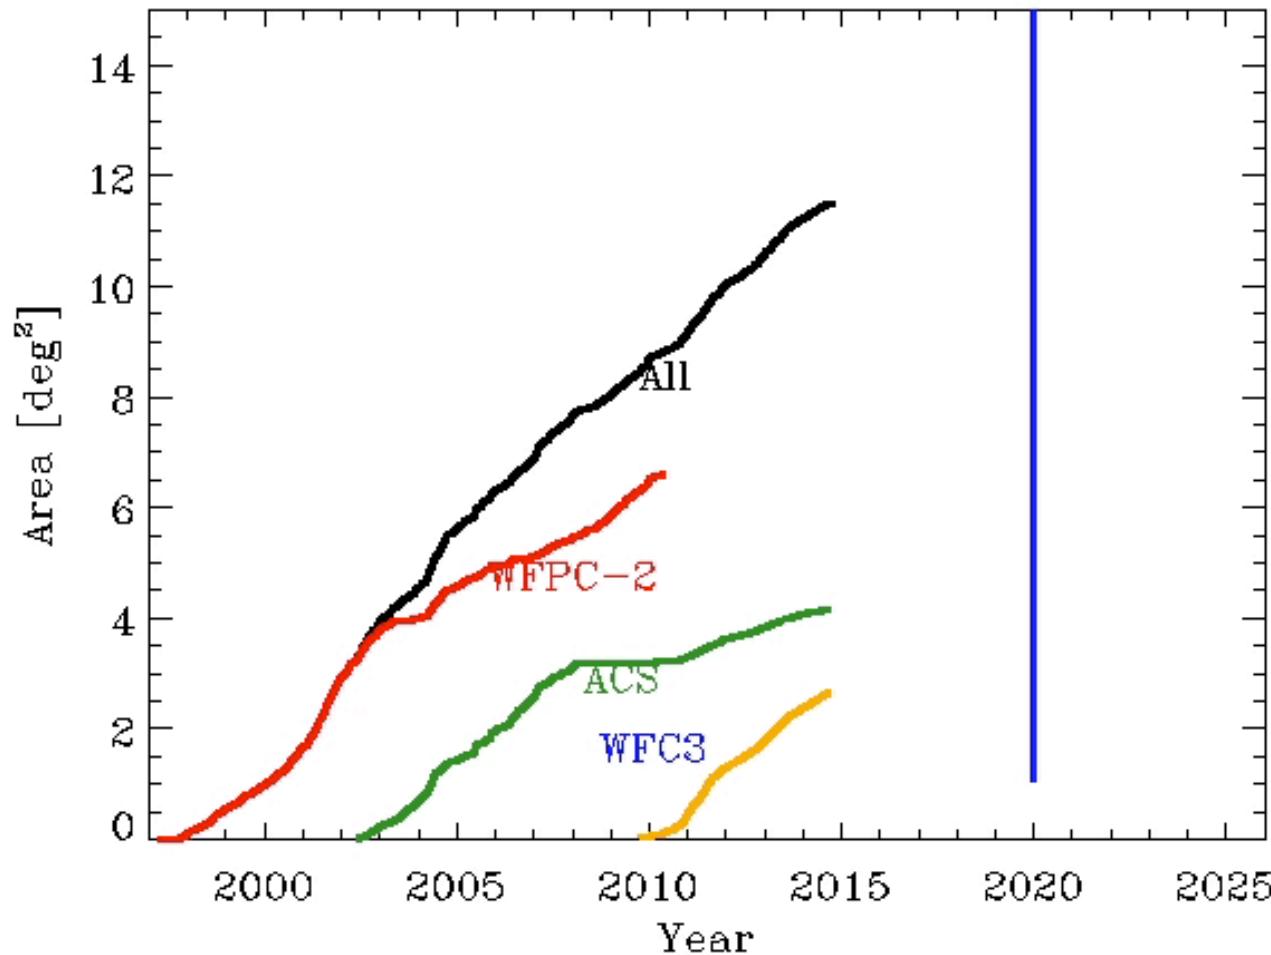
AS IN MANY OTHER DISCIPLINES THE BIG-DATA REVOLUTION HAS ARRIVED TO ASTRONOMY TOO



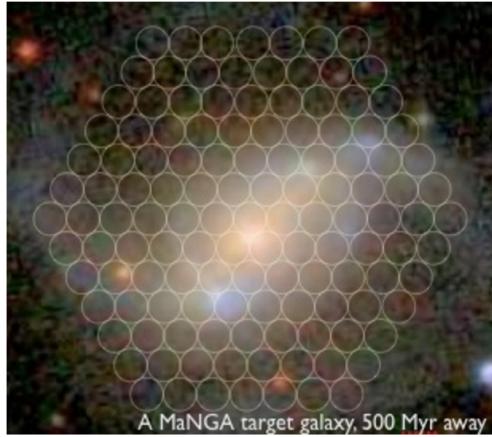
EXTREMELY LARGE
IMAGING SURVEYS
DELIVERING BILLIONS
OF OBJECTS IN 2-5 YEARS



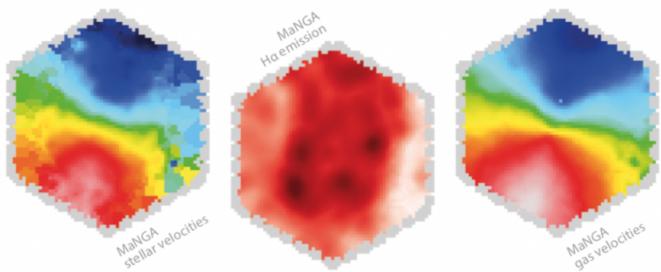
LSST simulation



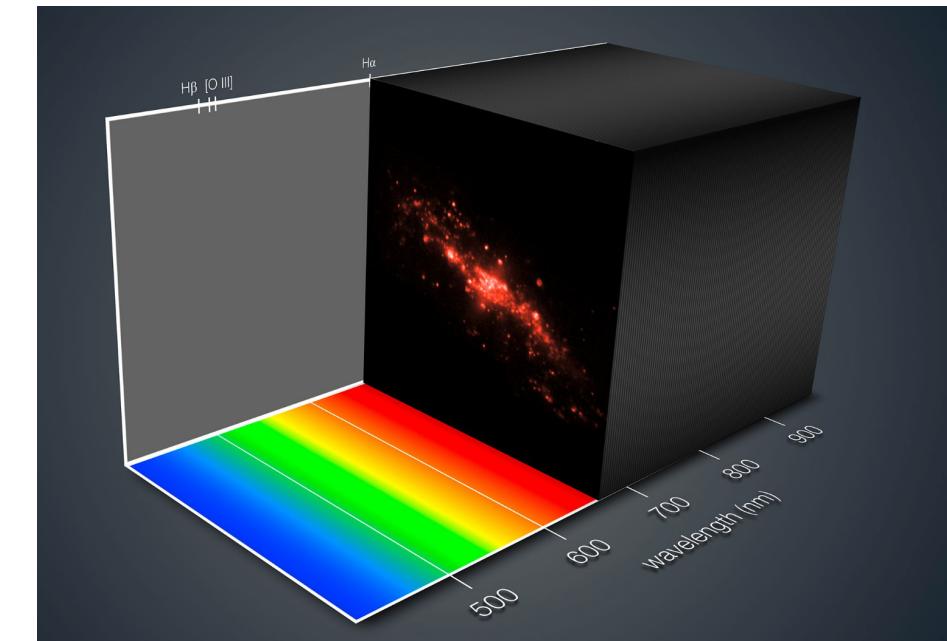
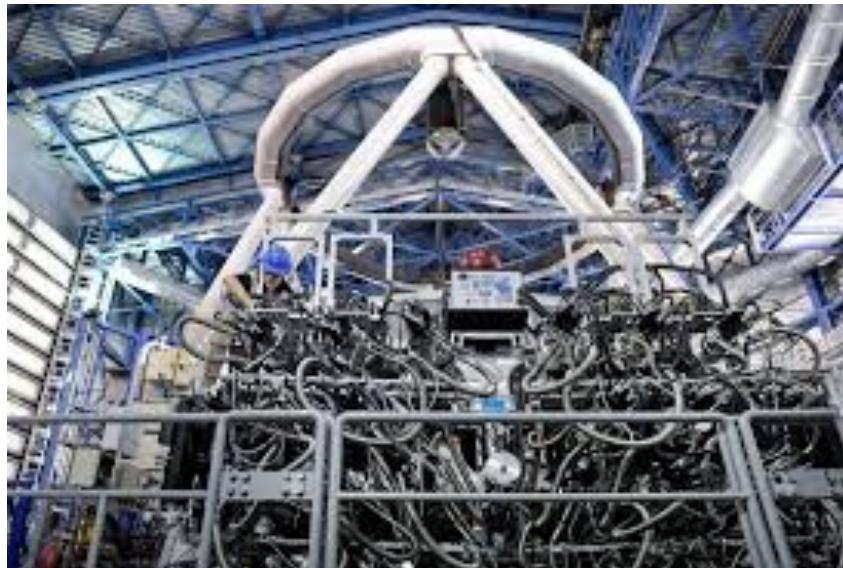
(Thanks to J. Brinchmann)



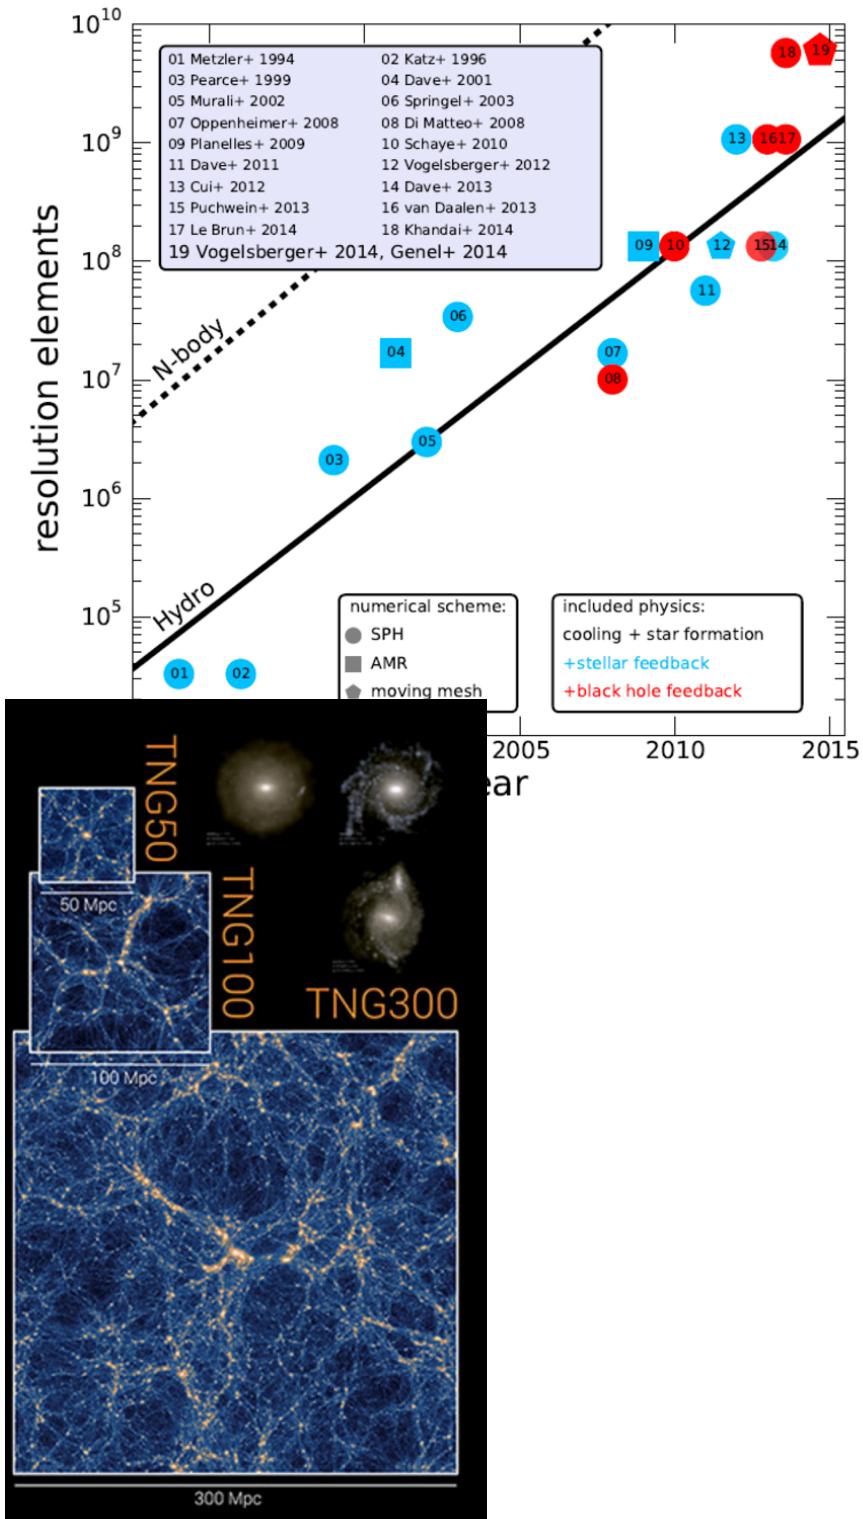
NOT ONLY VOLUME: AN
INCREASING
COMPLEXITY OF DATA



MANGA Survey

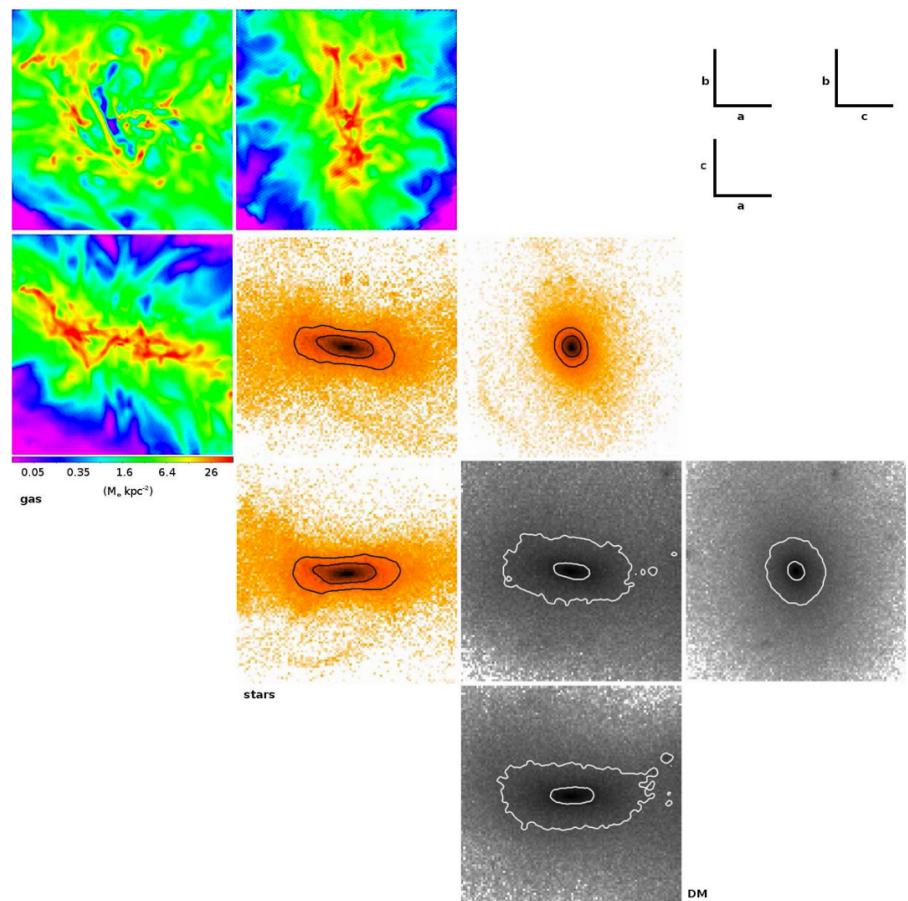


MUSE@VLT



Genel+14

AND ALSO
SIMULATIONS!



Ceverino+15

PROGRAM FOR THE TWO DAYS

- PART I: A VERY QUICK INTRODUCTION TO
‘CLASSICAL’ MACHINE LEARNING
 - UNSUPERVISED / SUPERVISED
 - GENERAL STEPS TO “TEACH A MACHINE”
 - “CLASSICAL” CLASSIFIERS

PROGRAM FOR THE TWO DAYS

- PART II: FOCUS ON ‘SHALLOW’ NEURAL NETWORKS
 - PERCEPTRON, NEURON DEFINITION
 - LAYER OF NEURONS, HIDDEN LAYERS
 - ACTIVATION FUNCTIONS
 - OPTIMIZATION [GRADIENT DESCENT, LEARNING RATES]
 - BACKPROPAGATION

PROGRAM FOR THE TWO DAYS

- **PART III: CONVOLUTIONAL NEURAL NETWORKS**
 - CONVOLUTIONS AS NEURONS
 - CNNs [POOLING, DROPOUT]
 - VANISHING GRADIENT / BATCH NORMALIZATION

PROGRAM FOR THE TWO DAYS

- PART IV: IMAGE TO IMAGE NETOWRKS +
INTRODUCTION TO UNSUPERVISED DEEP LEARNING
 - NETWORKS FOR IMAGE SEGMENTATION
 - AUTO-ENCODERS
 - GENERATIVE ADVERSARIAL NETOWRKS
 - ANOMALY DETECTION

PROGRAM FOR THE TWO DAYS

- **PART V: SOME PRACTICAL CONSIDERATIONS**
 - HOW DO I SETUP MY CNN?
 - HOW LARGE DO TRAINING SETS NEED TO BE?
 - OPTIMIZING YOUR NET: HYPER PARAMETER SEARCH
 - VISUALIZING CNNs [DECONVNETS,
INCEPTIONISM, INTEGRATED GRADIENTS]

HANDS-ON SESSION

WE WILL TRY TO IMPLEMENT SOME OF THE THINGS
LEARNED

WE WILL TEST SEVERAL APPROACHES TO CLASSIFY
GALAXIES

LET'S TRY TO DISCUSS AS MUCH AS POSSIBLE!

SOFTWARE REQUIREMENTS

- PYTHON 3 OR GREATER
- TENSORFLOW FOR DEEP LEARNING
- KERAS - HIGH LEVEL LIBRARY WHICH MAKES GPU CODING TRANSPARENT - SIMPLIFIES THINGS A LOT AND MOST OF THE TIME ENOUGH FOR OUR APPLICATIONS

PART I: AN INTRODUCTION TO
“CLASSICAL” MACHINE LEARNING

WHAT DOES MACHINE LEARNING DO?

the machine is told what to look for

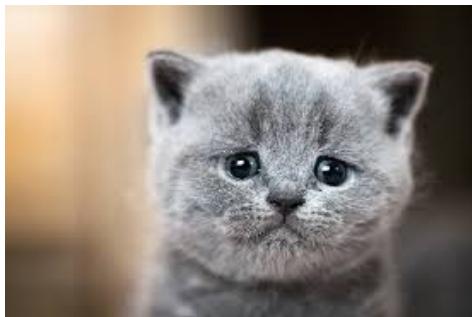
SUPERVISED

the machine is NOT told what to look for

UN-SUPERVISED

**TWO BIG TYPES OF MACHINE LEARNING
ALGORITHMS**

CAT



CAT

SUPERVISED LEARNING

DOG



HUMAN LABELLING

DOG



the machine is told what to look for

CAT



CAT



SUPERVISED LEARNING

DOG



DOG



the machine is told what to look for

HUMAN LABELLING

TRAINING SET
OF LABELED
EXAMPLES

CAT



CAT



DOG



DOG



SUPERVISED LEARNING



ML



CAT

CAT



CAT



DOG



DOG



SUPERVISED LEARNING



ML



DOG

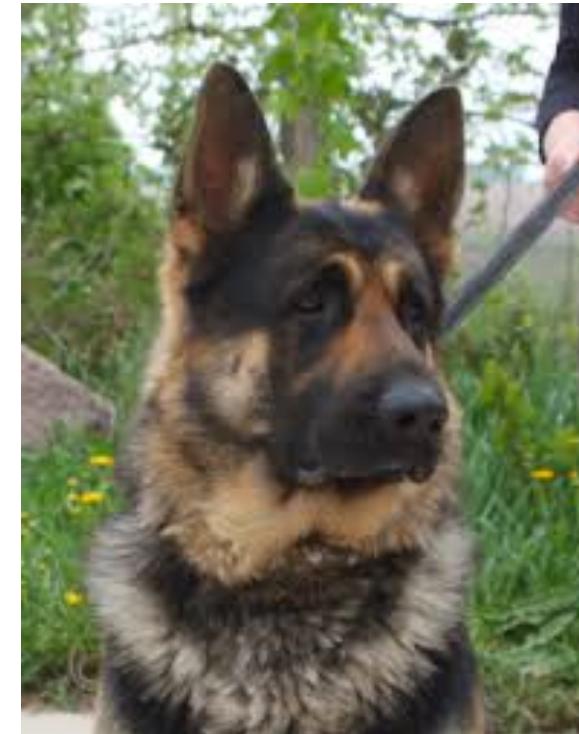
UNSUPERVISED LEARNING



UNSUPERVISED LEARNING



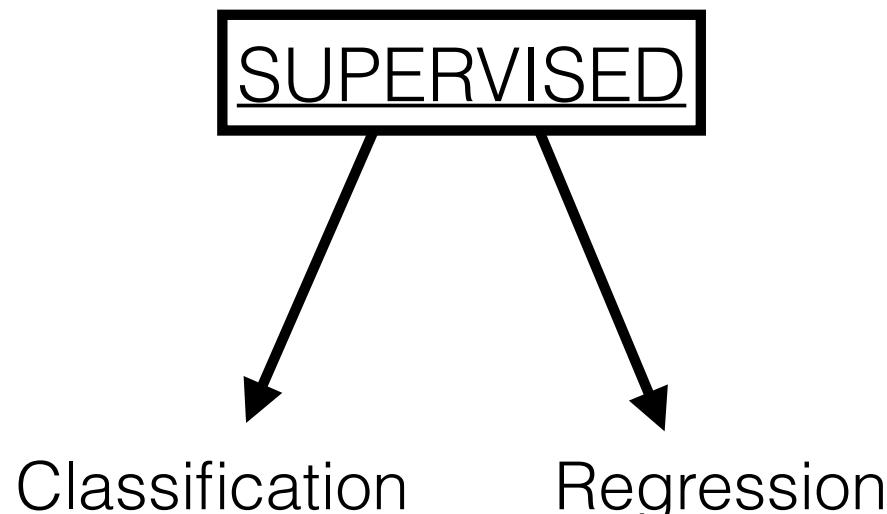
UNSUPERVISED LEARNING



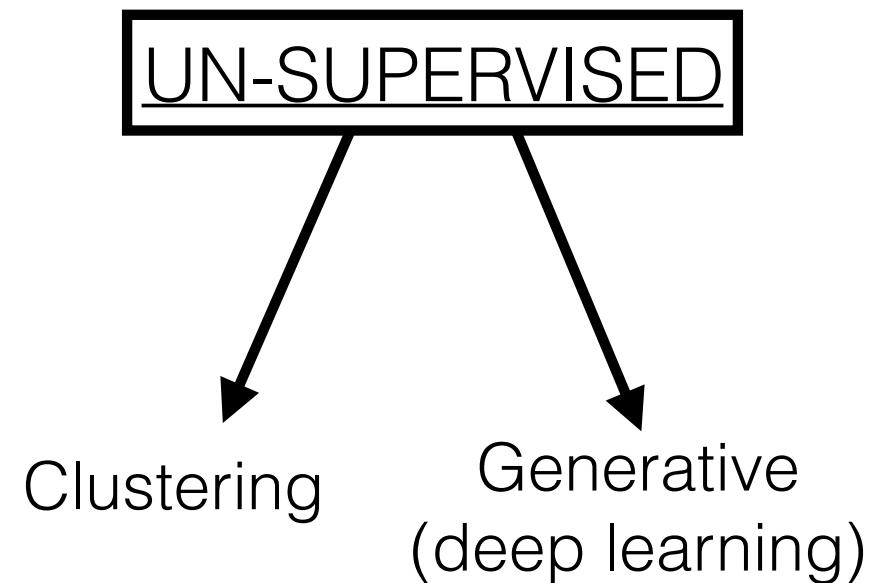
THE DEFINITION OF CLASSES IS SOMETIMES
NOT OBVIOUS

WHAT DOES MACHINE LEARNING DO?

the machine is told what to look for

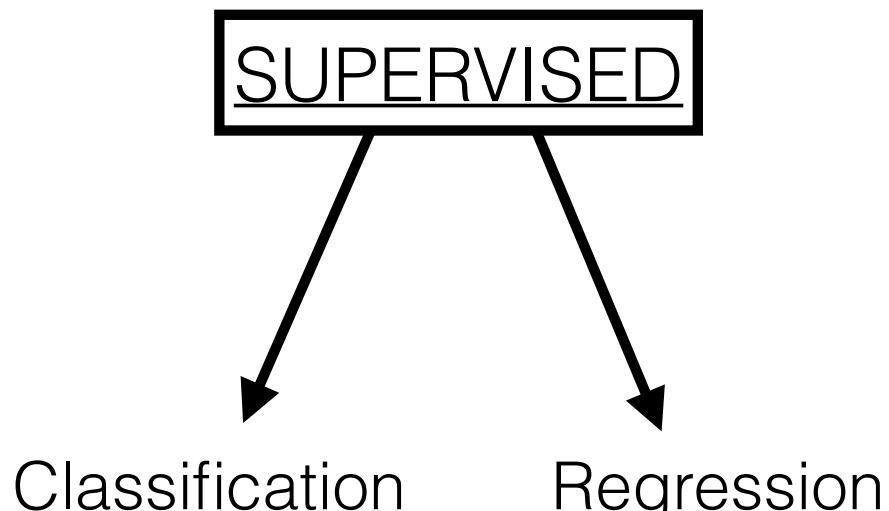


the machine is NOT told what to look for

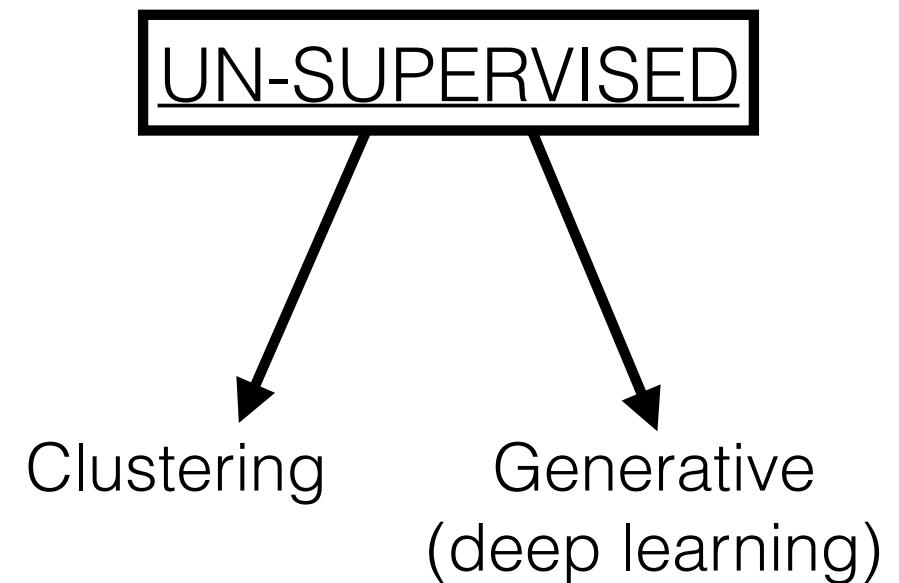


WHAT DOES MACHINE LEARNING DO?

the machine is told what to look for



the machine is NOT told what to look for



WHAT DOES MACHINE LEARNING DO?

SUPERVISED

Classification

Regression

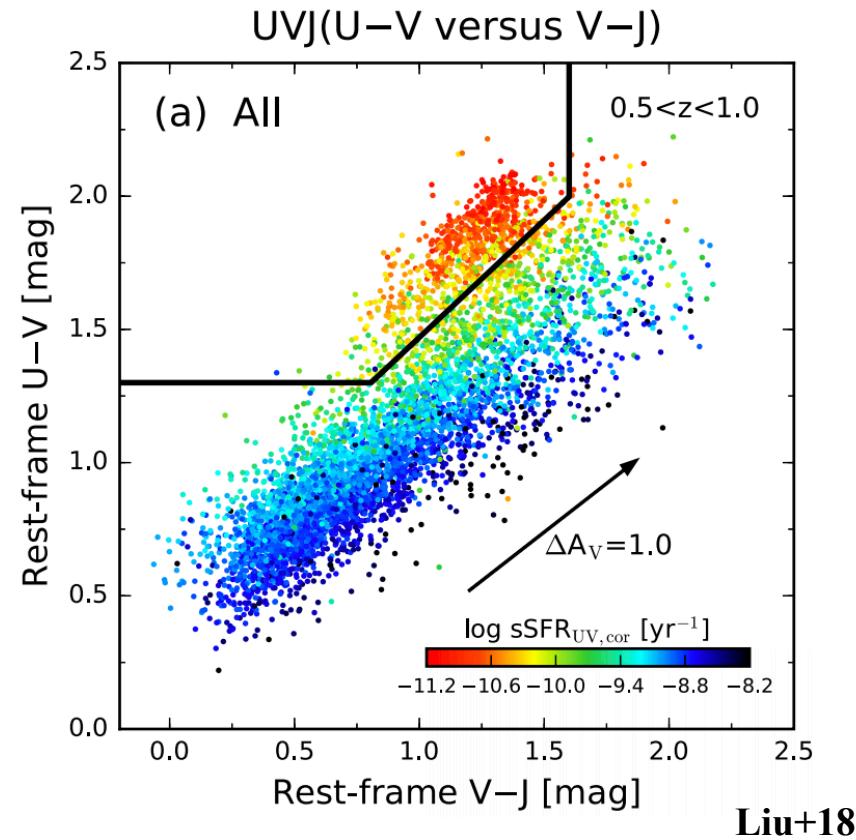
UN-SUPERVISED

Clustering

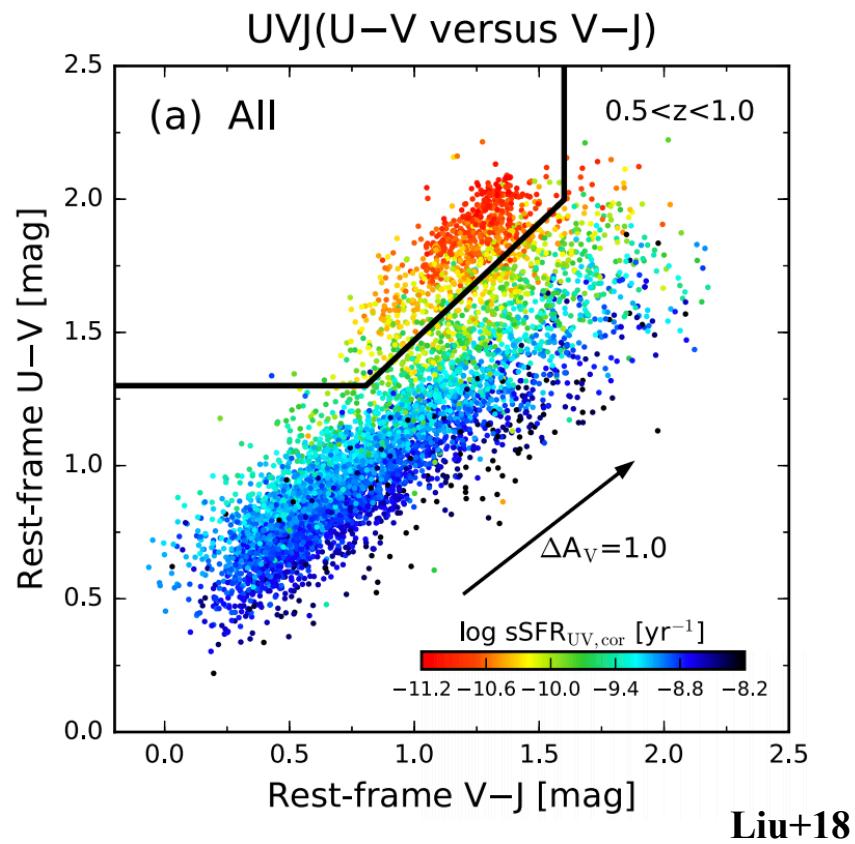
Generative
(deep learning)

DEEP LEARNING

THRE IS NO MAGIC IN MACHINE LEARNING, AND IT IS ACTUALLY PRETTY SIMPLE

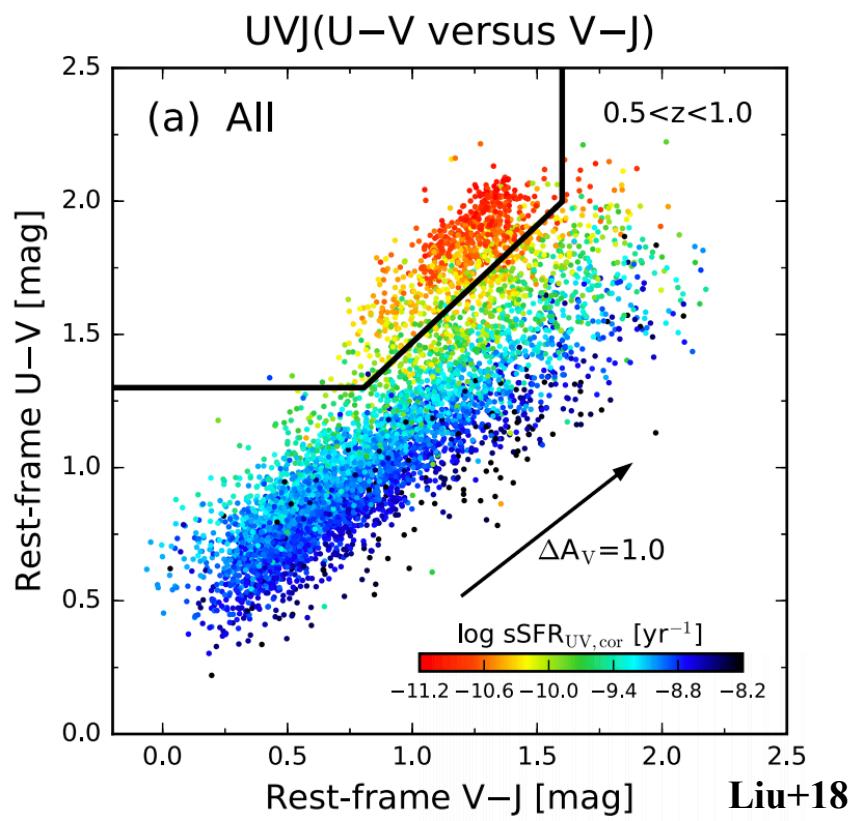


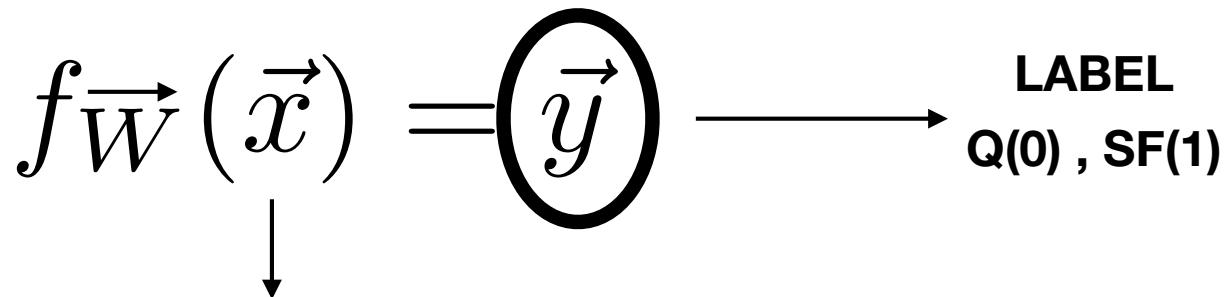
$$f_W(\vec{x}) = \vec{y}$$



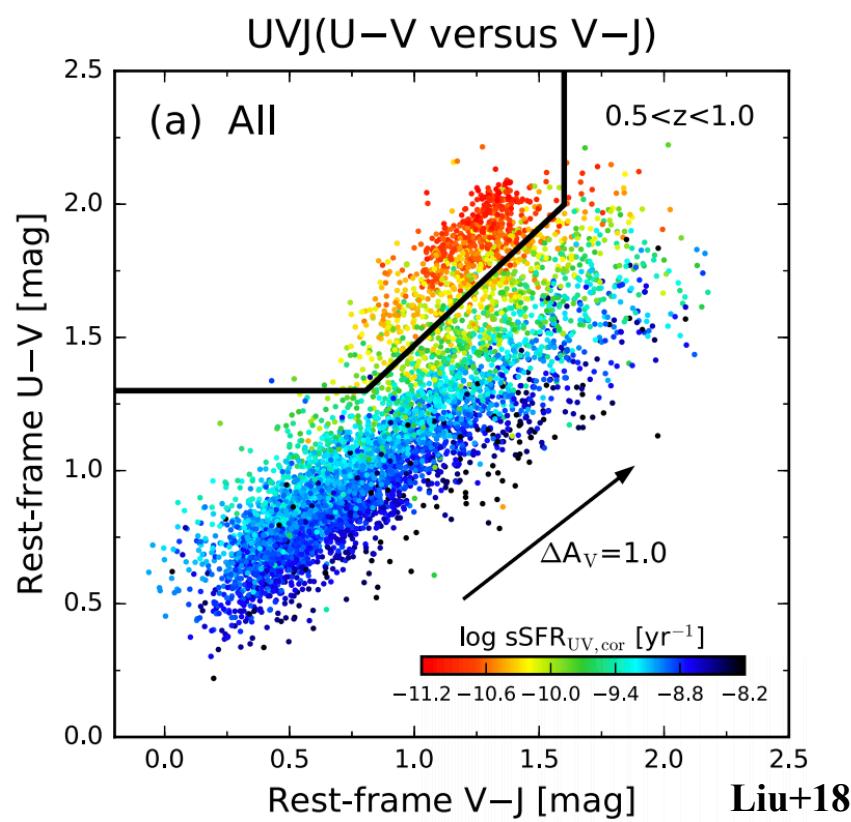
$$f_W(\vec{x}) = \vec{y}$$

LABEL
Q , SF





(U-V, V-J) FEATURES



$$f_{\vec{W}}(\vec{x}) = \vec{y} \longrightarrow \text{LABEL}$$

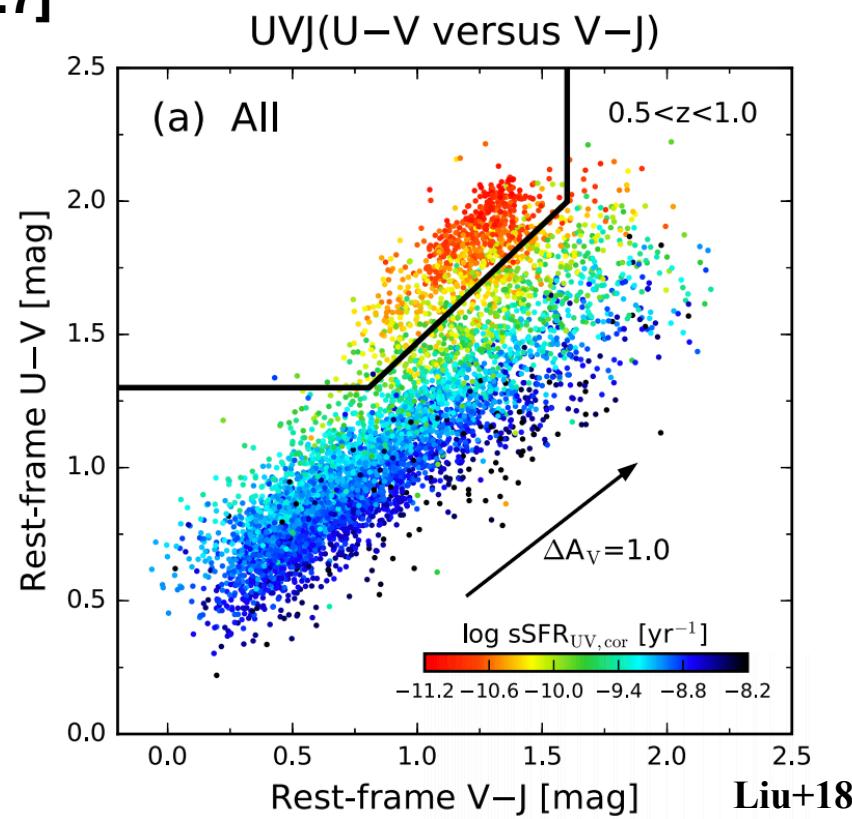
Q(0) , SF(1)

NETWORK FUNCTION

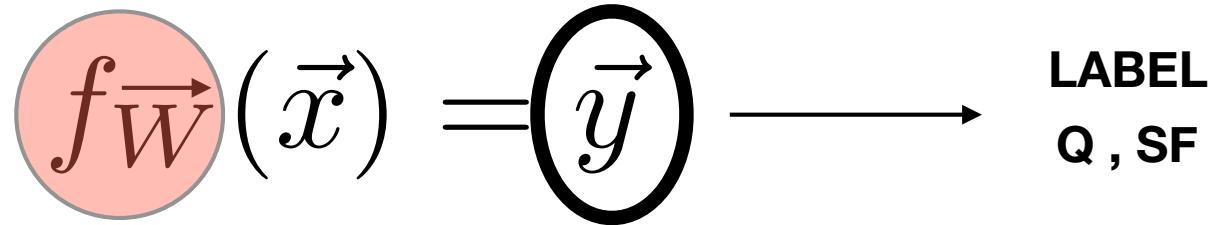
(U-V, V-J) FEATURES

$$\text{sgn}[(u-v)-0.8*(v-j)-0.7]$$

WEIGHTS



**“CLASSICAL”
MACHINE LEARNING**



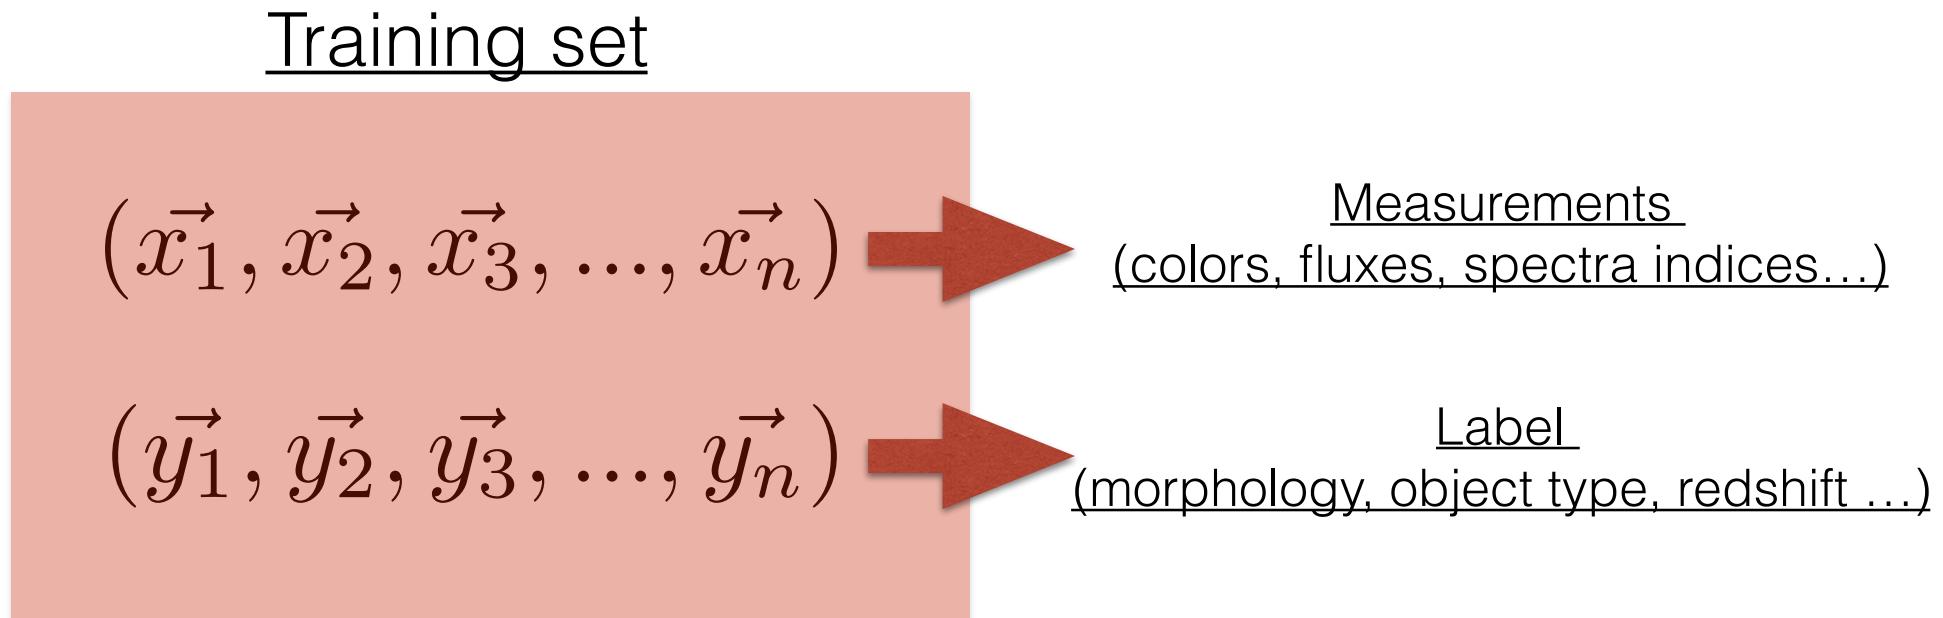
$$\text{sgn}[(u-v)-W_1*(v-j)-W_2]$$



**REPLACE THIS BY A GENERAL
NON LINEAR FUNCTION WITH SOME PARAMETERS W**

SUPERVISED LEARNING

Given a dataset with known labels (measurements) - find a function that can assign (predict) measurements for an unlabeled dataset



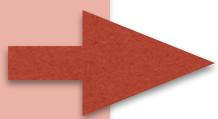
SUPERVISED LEARNING

Given a dataset with known labels (measurements) - find a function that can assign (predict) measurements for an unlabeled dataset

Training set

$$(\vec{x}_1, \vec{x}_2, \vec{x}_3, \dots, \vec{x}_n)$$

$$(\vec{y}_1, \vec{y}_2, \vec{y}_3, \dots, \vec{y}_n)$$

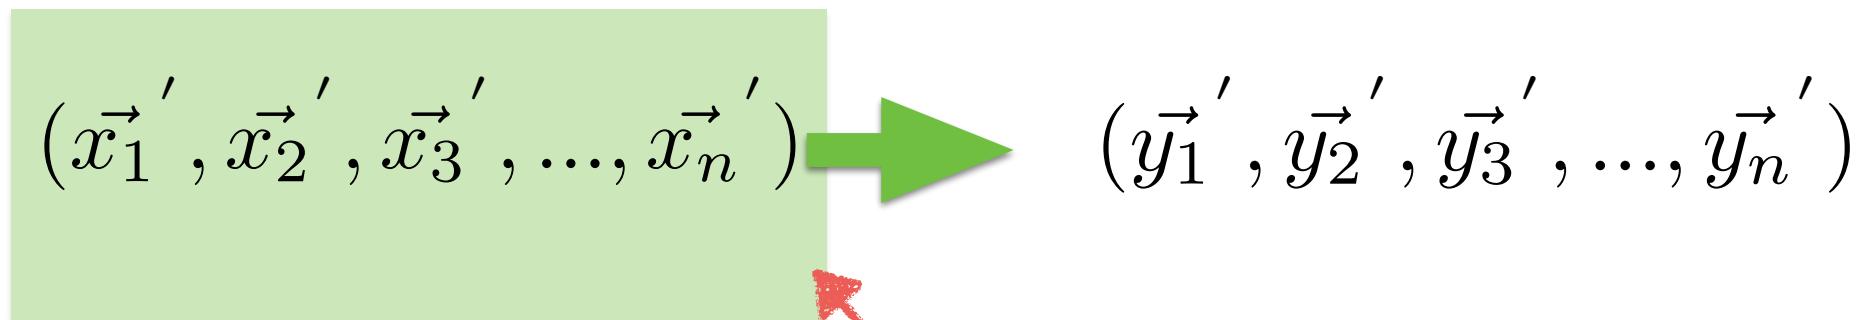


$$f_W(\vec{x}) = \vec{y}$$

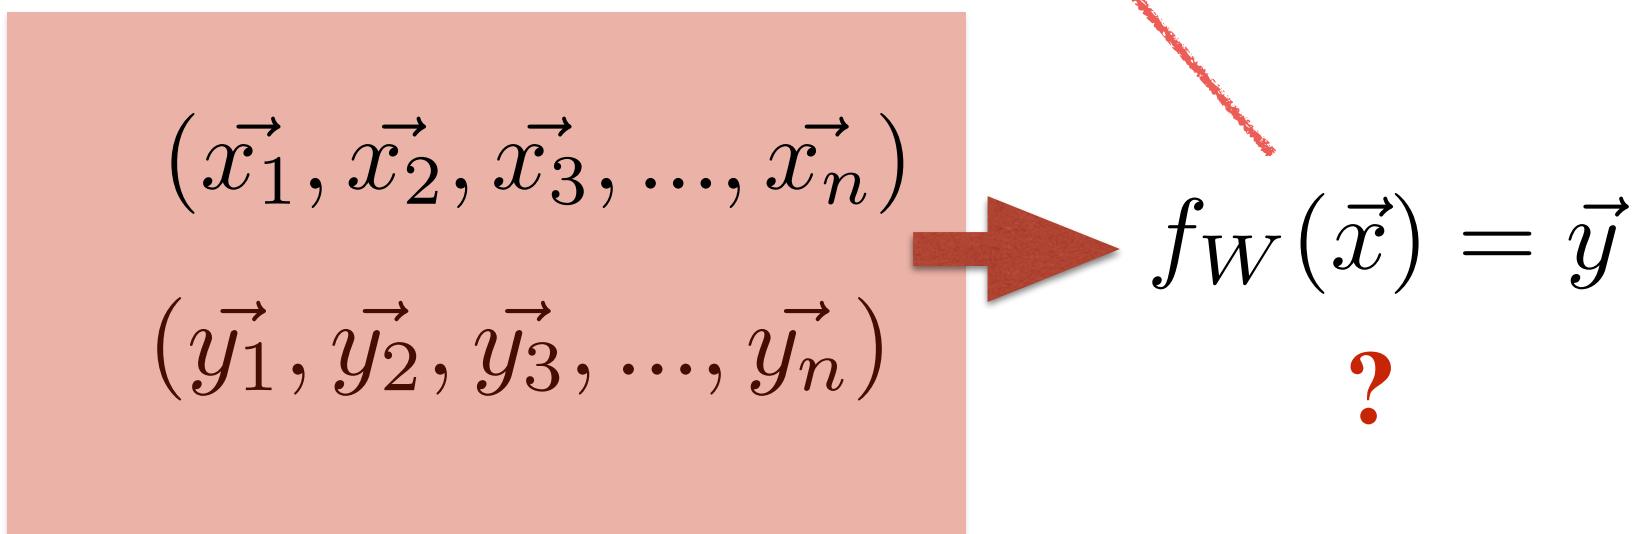
?

SUPERVISED LEARNING

Unlabeled set



Training set



$$(\vec{x}_1, \vec{x}_2, \vec{x}_3, \dots, \vec{x}_n)$$

$$\vec{x} \in \mathbb{R}^d$$

$$(\vec{y}_1, \vec{y}_2, \vec{y}_3, \dots, \vec{y}_n)$$

$$\vec{y} \in \mathbb{R} \quad \vec{y} \in \mathbb{N}$$

GENERAL GOAL: Find a (non-linear) function that outputs the correct class / measurement for a given input object:

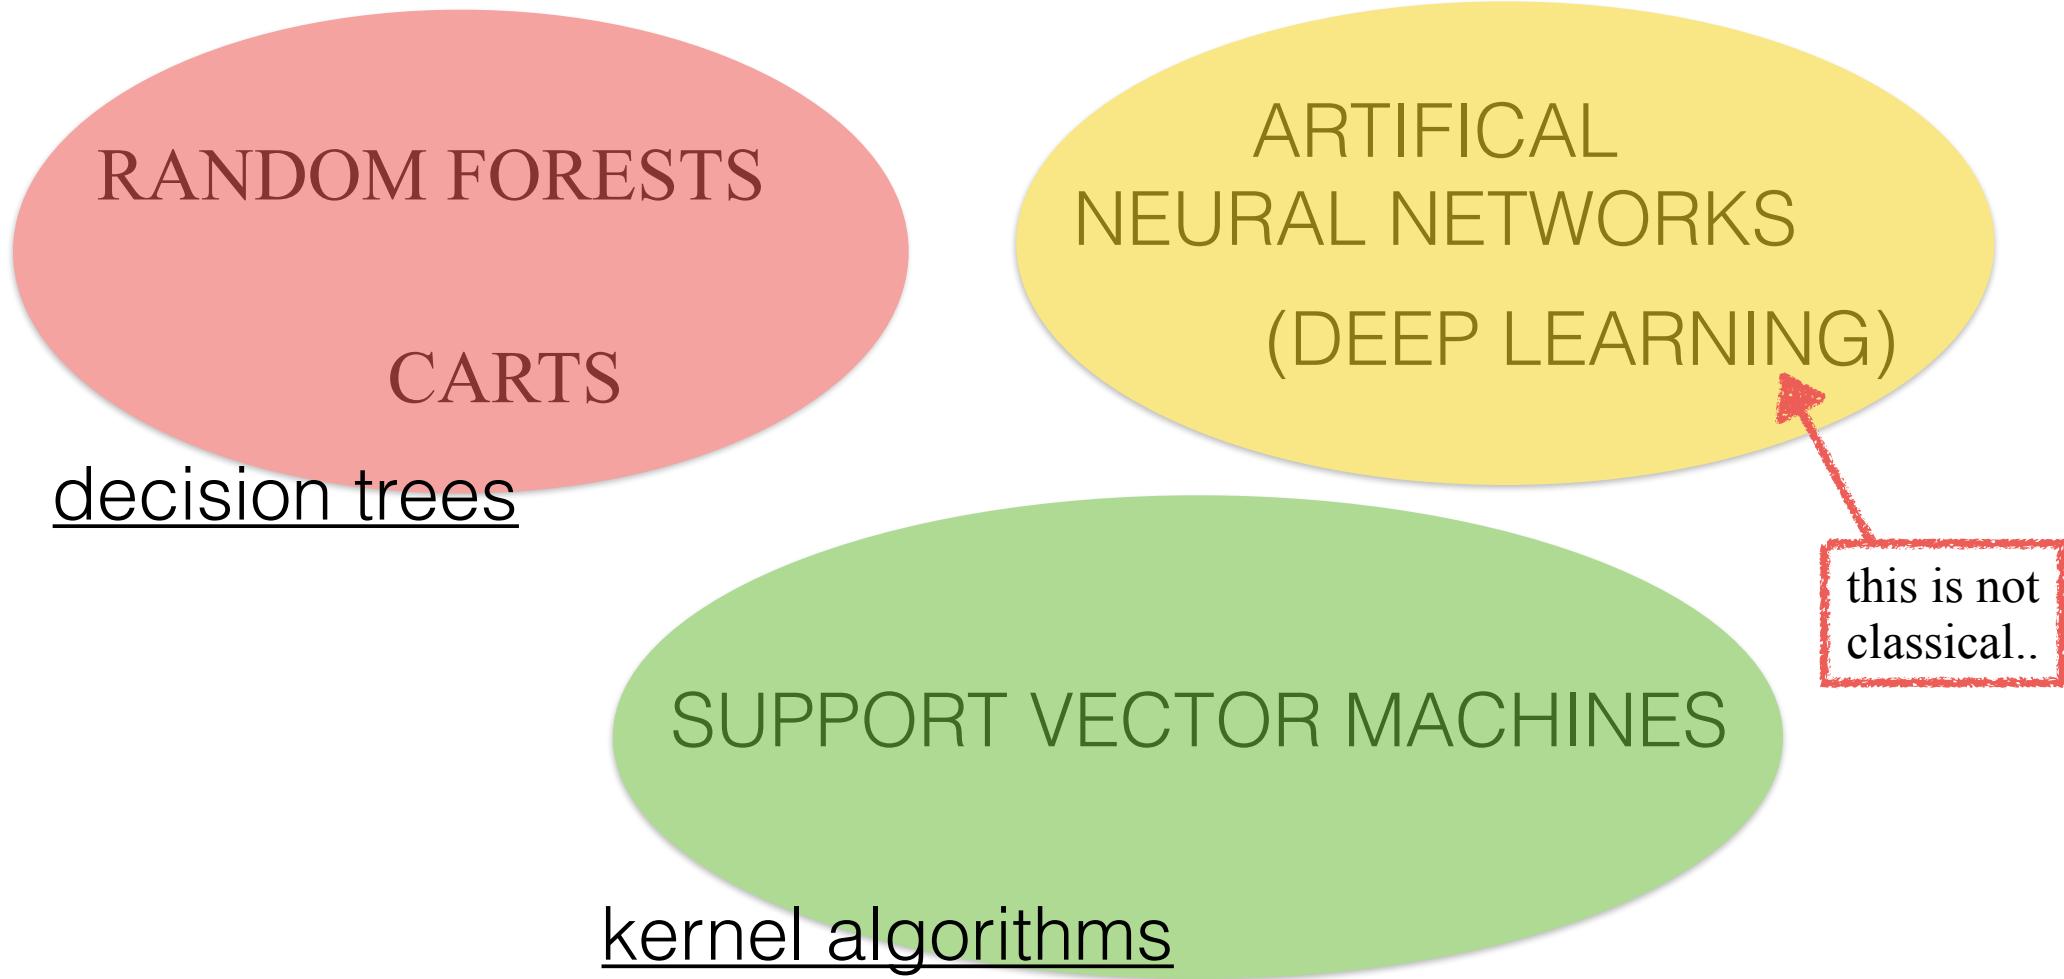
$$f_W(\vec{x})$$



Number of parameters - can be large

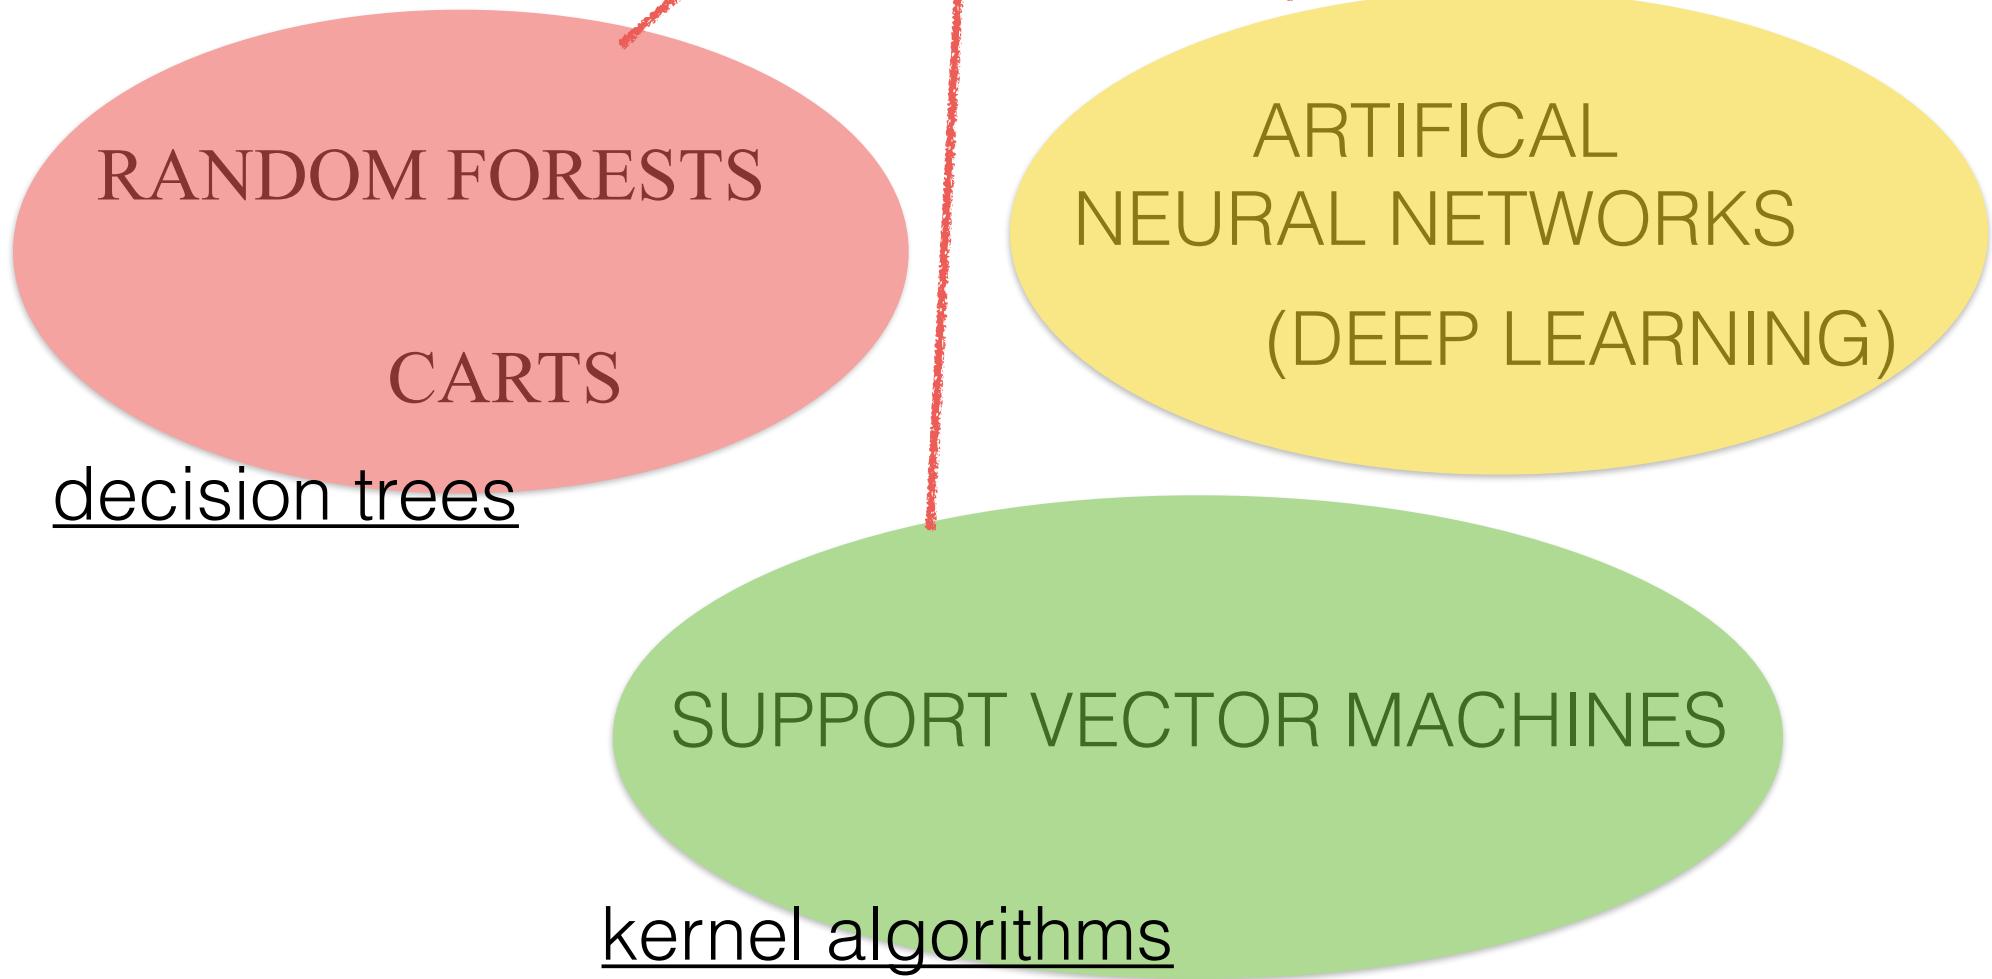
It is translated into a minimization problem : find \mathbf{W} such as the prediction error is minimal over all unseen vectors

Different “classical” supervised machine learning methods



The differences are
in the function
that is used

$$f_W(\vec{x})$$



We need two key elements

1. A LOSS FUNCTION

**2. A MINIMIZATION OR OPTIMIZATION
ALGORITHM**

We need two key elements

1. A LOSS FUNCTION

**2. A MINIMIZATION OR OPTIMIZATION
ALGORITHM**

**THIS IS COMMON TO ALL MACHINE LEARNING
ALGORITHMS**

1. DEFINE A LOSS FUNCTION

$$loss(F_W(.), \vec{x}_i, \vec{y}_i)$$

For example: $(F_W(\vec{x}_i) - \vec{y}_i)^2$ Quadratic loss function

2. MINIMIZE THE EMPIRICAL RISK

$$\mathfrak{R}_{empirical}(W) = \frac{1}{N} \sum_i^N [loss(W, \vec{x}, \vec{y})]$$



MINIMIZE THE RISK

EMPIRICAL RISK?

$$\mathfrak{R}_{\text{empirical}}(W) = \frac{1}{N} \sum_i^N [\text{loss}(W, \vec{x}, \vec{y})]$$

WE ARE MINIMIZING WITH RESPECT TO A FINITE NUMBER OF OBSERVED EXAMPLES

EMPIRICAL RISK?

$$\mathfrak{R}_{\text{empirical}}(W) = \frac{1}{N} \sum_i^N [\text{loss}(W, \vec{x}, \vec{y})]$$

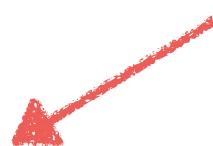
WE ARE MINIMIZING WITH RESPECT TO A FINITE NUMBER OF OBSERVED EXAMPLES

OBSERVED DATASET



EMPIRICAL RISK?

$$\mathfrak{R}_{\text{empirical}}(W) = \frac{1}{N} \sum_i^N [\text{loss}(W, \vec{x}, \vec{y})]$$



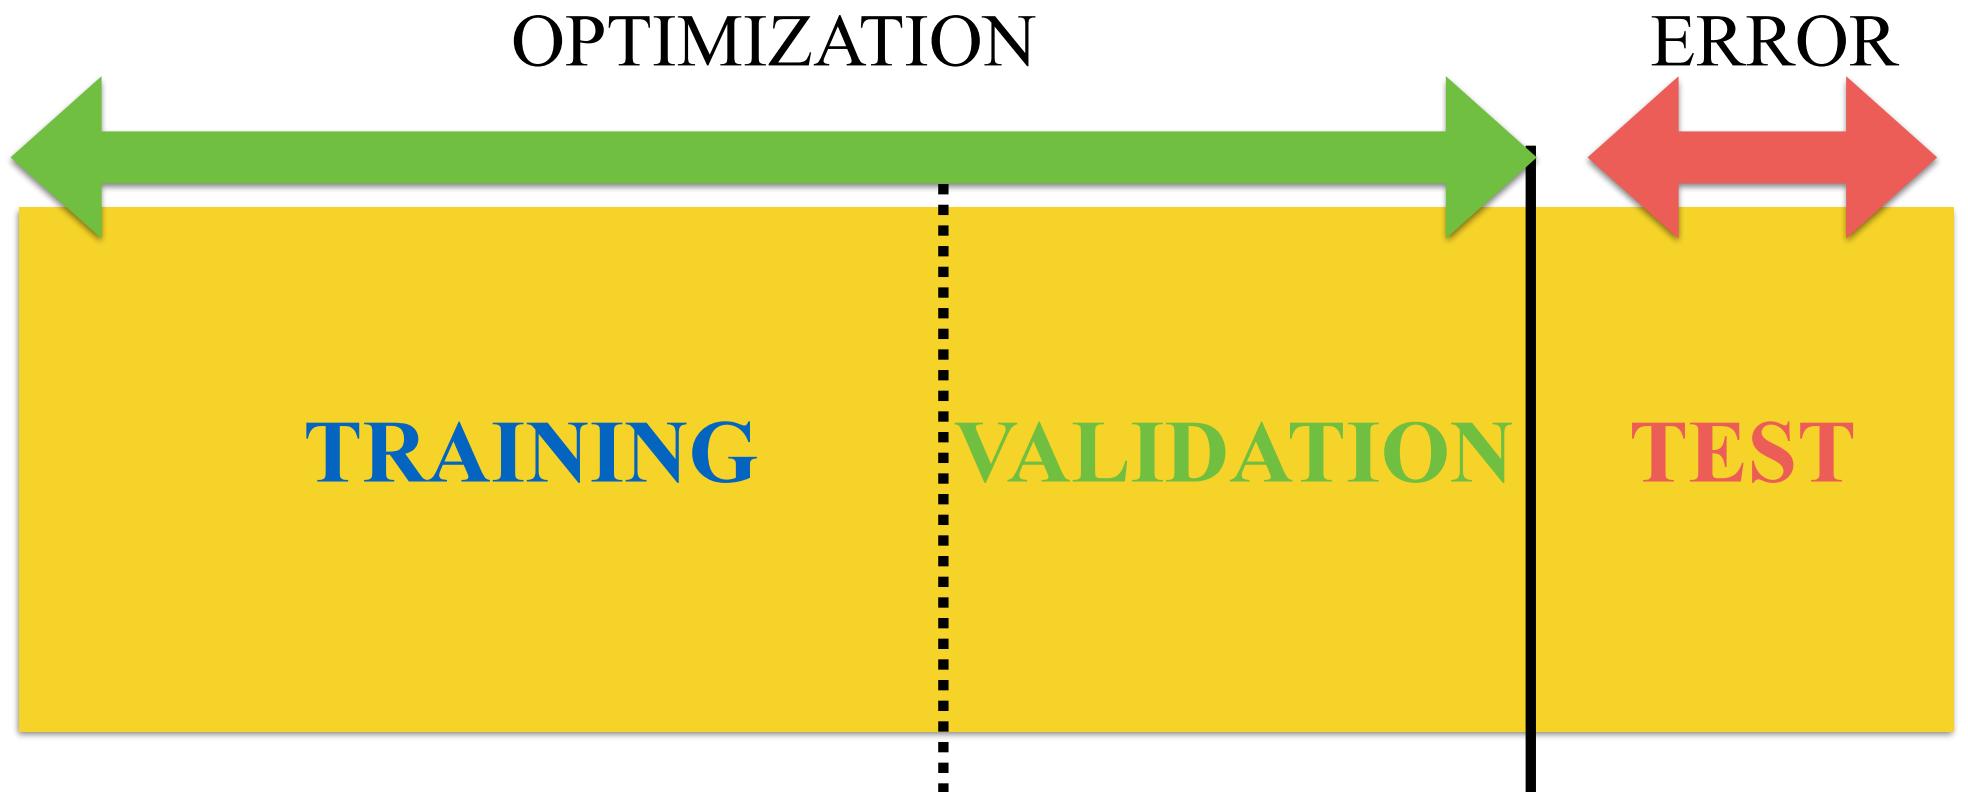
WE ARE MINIMIZING WITH RESPECT TO A FINITE NUMBER OF OBSERVED EXAMPLES

ALL “GALAXIES IN THE UNIVERSE”

OBSERVED DATASET



In practice

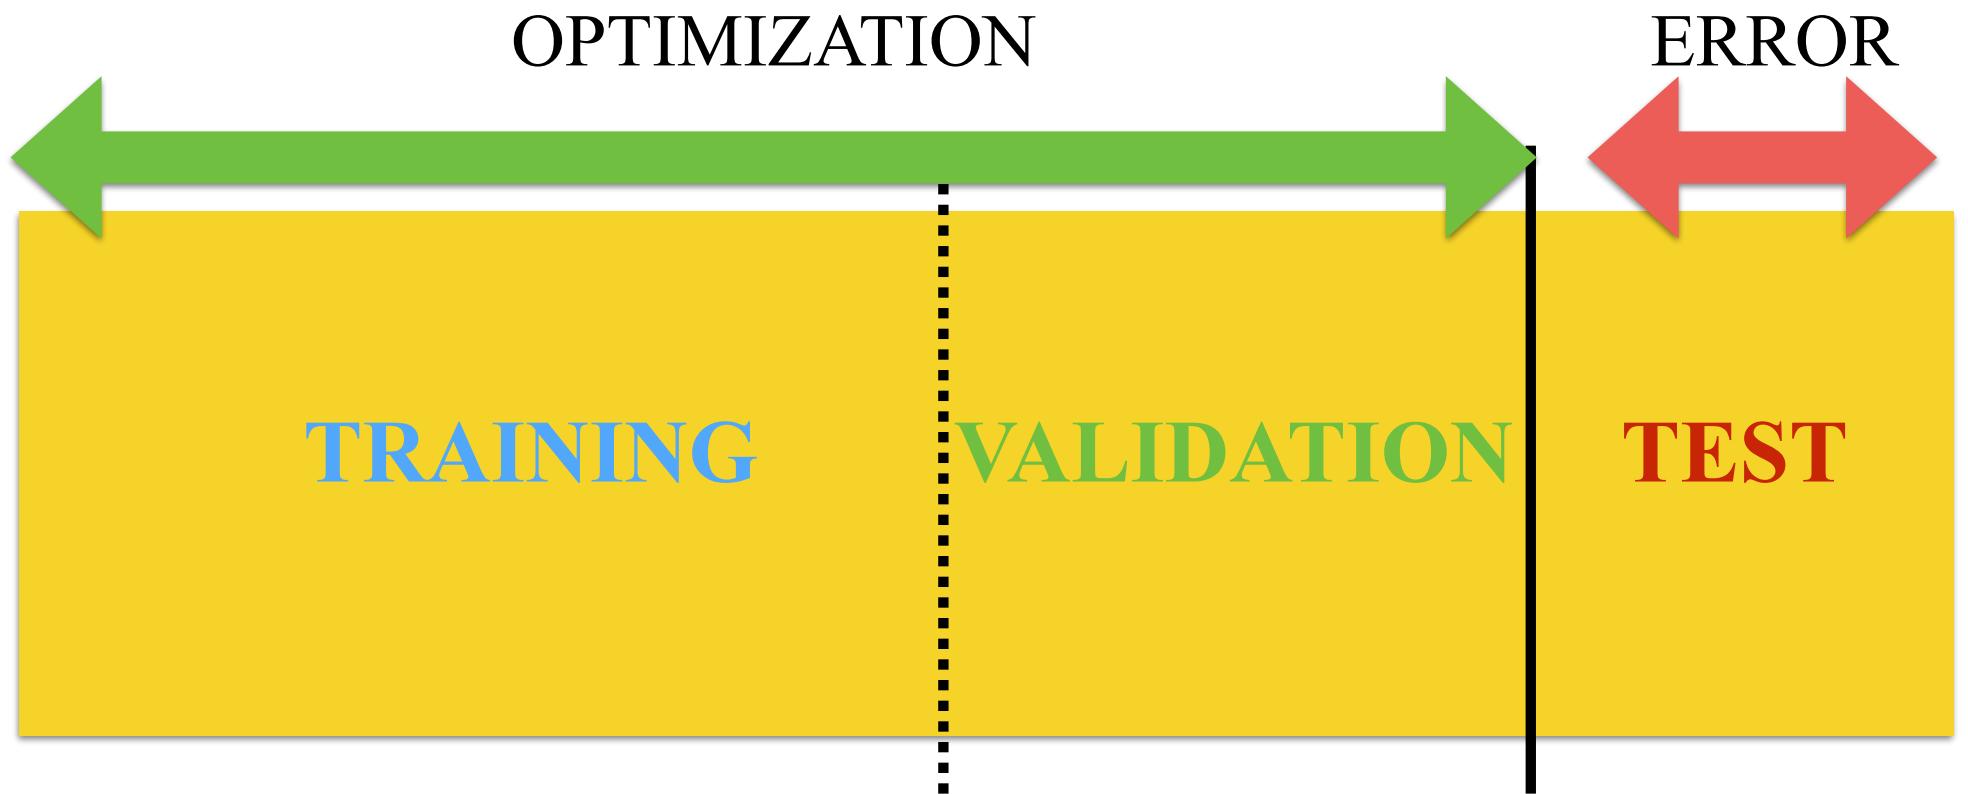


training set: use to train the classifier

validation set: use to monitor performance in real time - check
for overfitting

test set: use to train the classifier

In practice



**NO CHEATING! NEVER USE TRAINING TO VALIDATE
YOUR ALGORITHM!**

The algorithm used to minimize is
called OPTIMIZATION

THERE ARE SEVERAL OPTIMIZATION TECHNIQUES

Optimization

THERE ARE SEVERAL OPTIMIZATION TECHNIQUES

THEY DEPEND ON THE MACHINE LEARNING ALGORITHM

Optimization

THERE ARE SEVERAL OPTIMIZATION TECHNIQUES

THEY DEPEND ON THE MACHINE LEARNING ALGORITHM

NEURAL NETWORKS USE THE GRADIENT DESCENT AS WE
WILL SEE LATER

$$W_{t+1} = W_t - \lambda_h \nabla f(W_t)$$

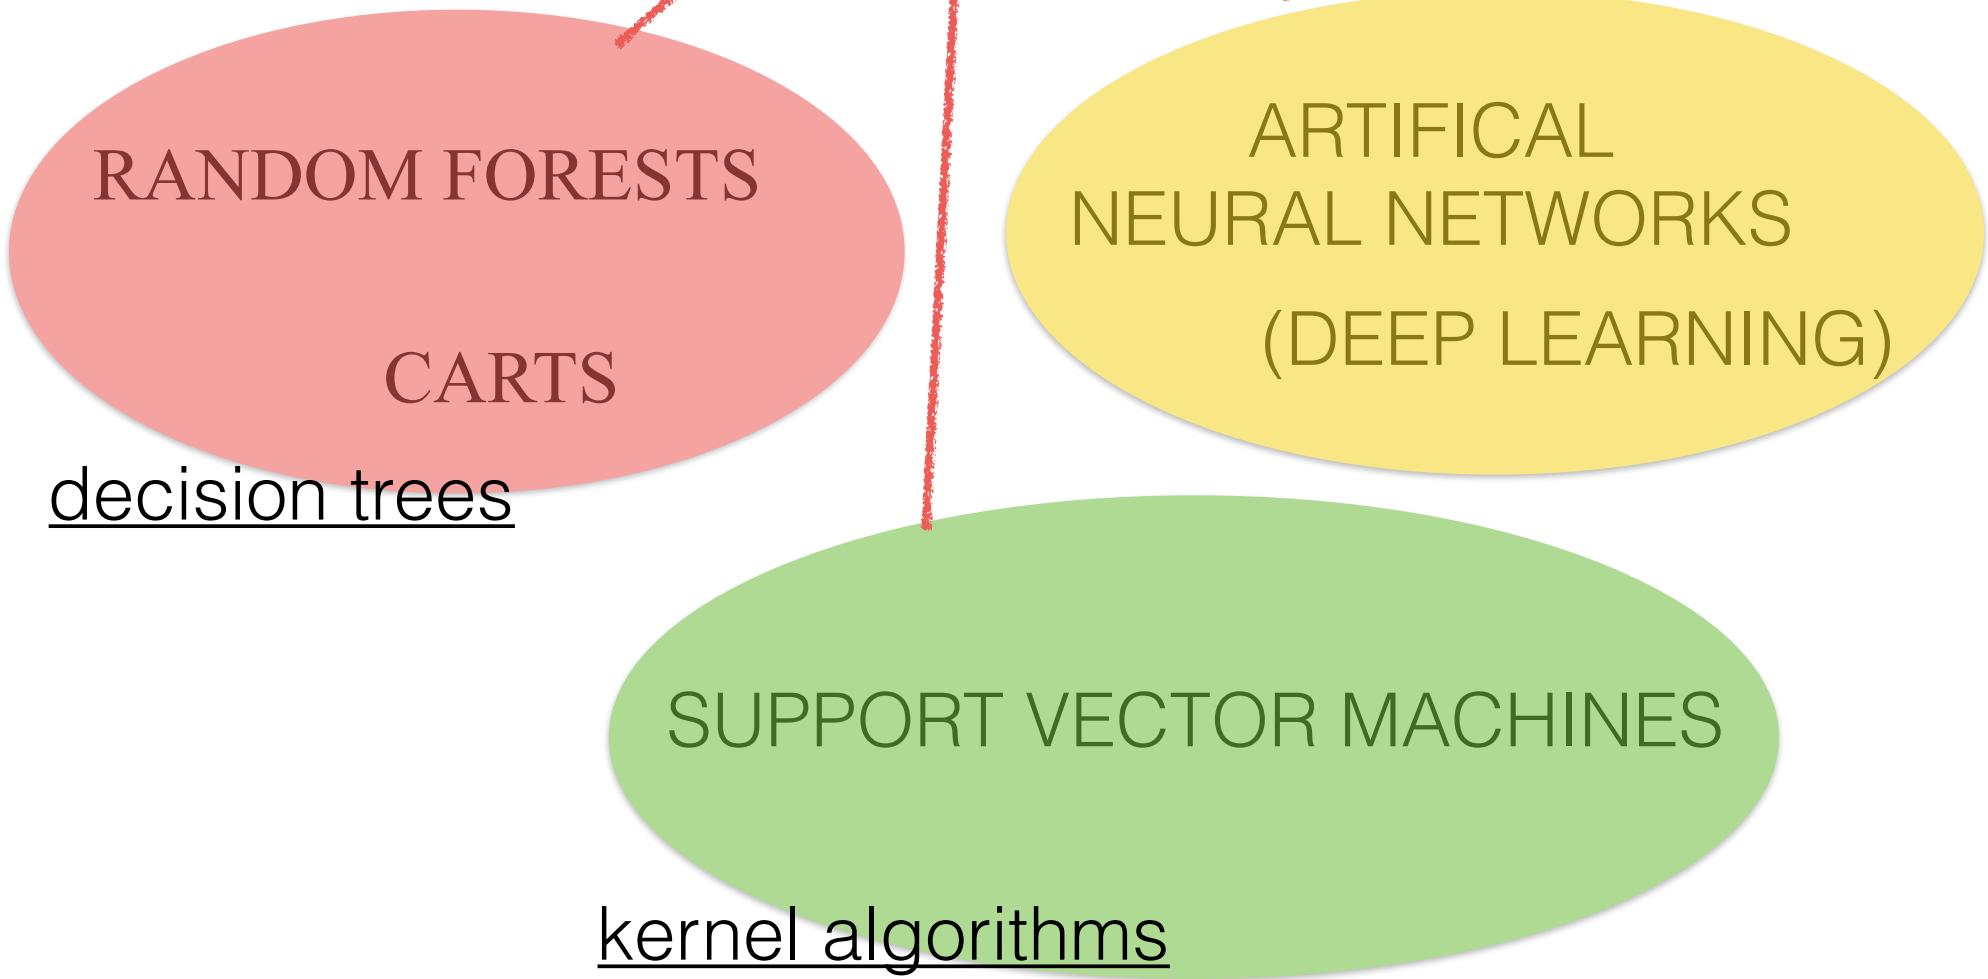
weights to be learned

epoch

learning rate

The differences are
in the function
that is used

$$f_W(\vec{x})$$



DECISION TREES ALGORITHMS

DECISION TREES

CARTS

Classification Trees

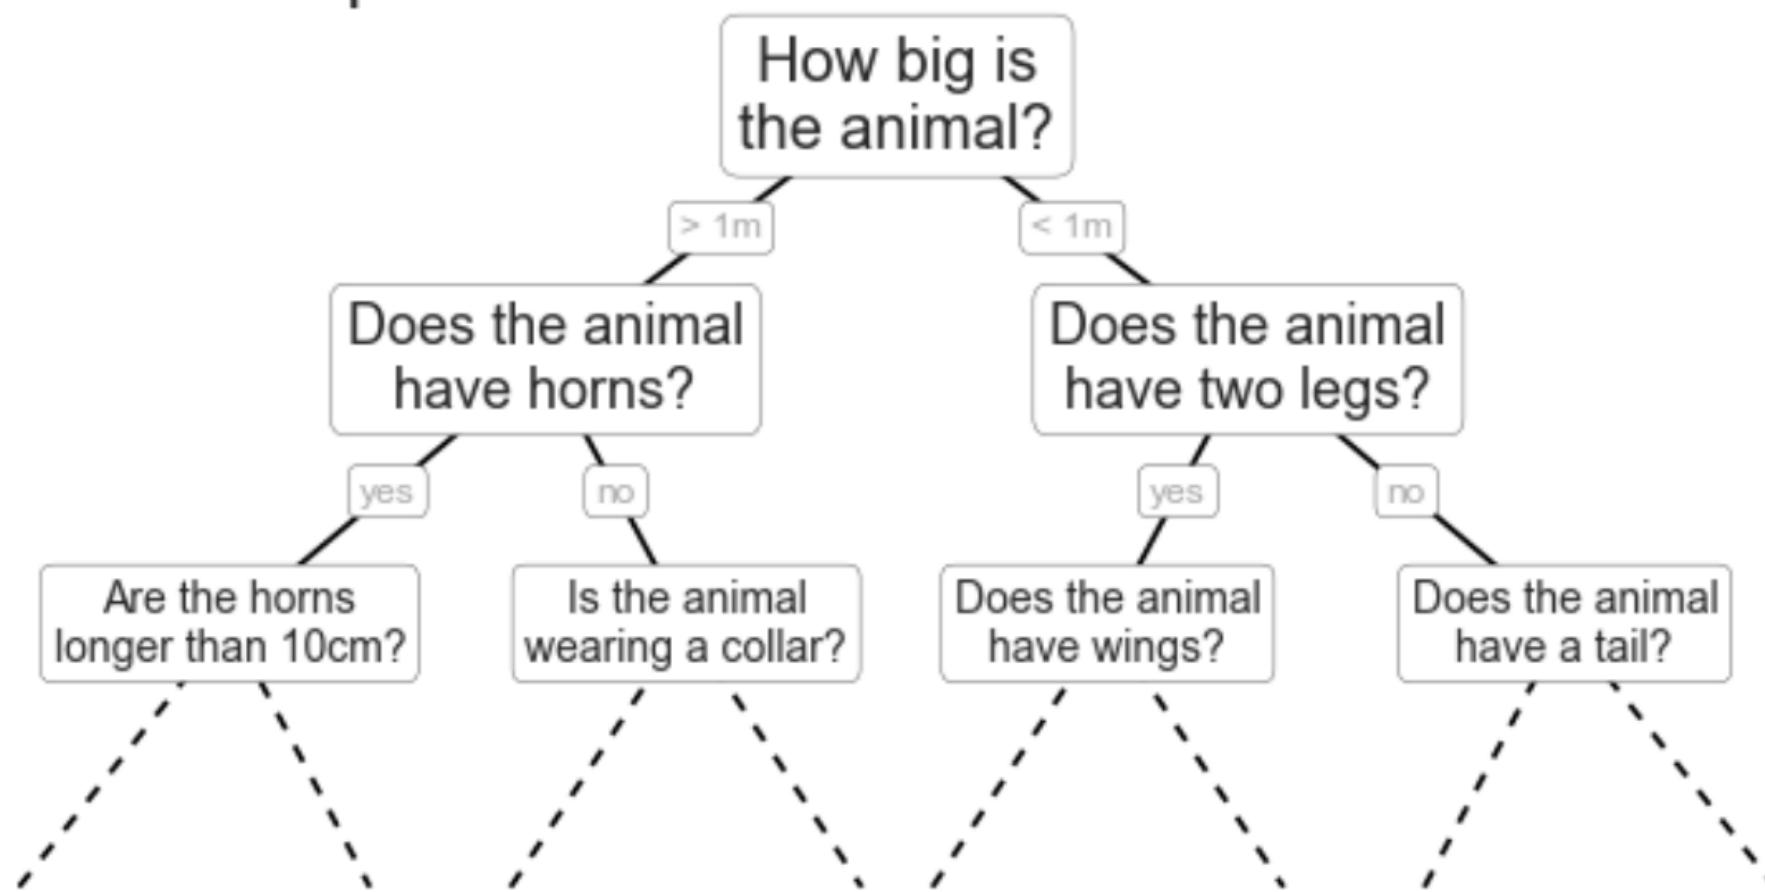
Regression Trees

BOOSTED TREES

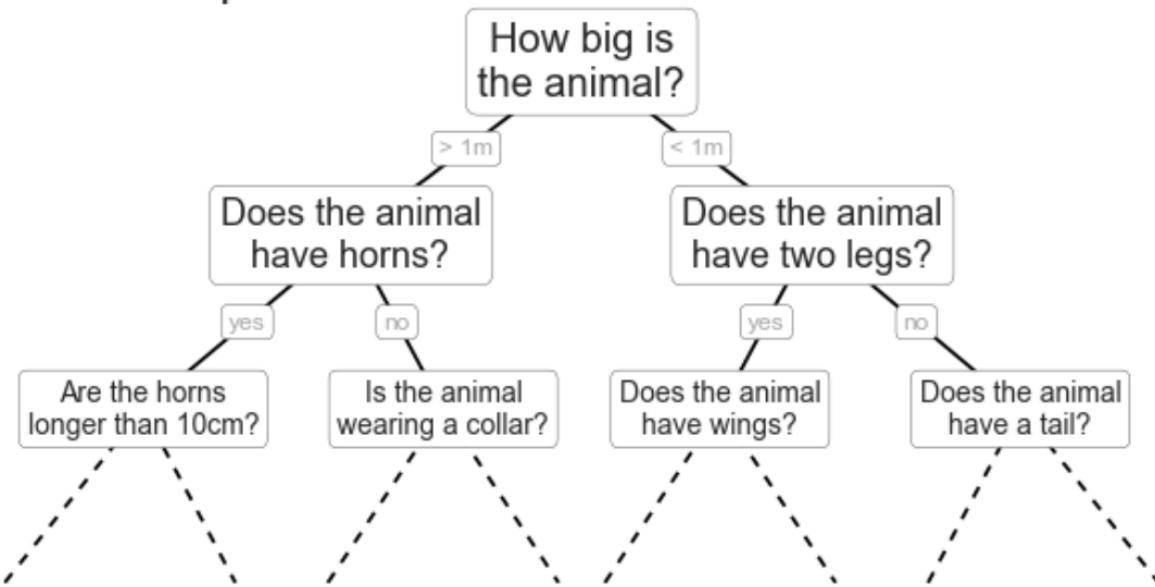
Random Forests

CLASSIFICATION AND REGRESSION TREES (CARTS)

THIS IS THE SIMPLEST AND MORE INTUITIVE MACHINE LEARNING ALGORITHM



DECISION TREES (CARTS)



IT IS BUILT IN AN ITERATIVE WAY

1. FROM THE INPUT PARAMETERS, FIRST FIND THE PROPERTY THAT BEST SPLITS INTO 2 GROUPS [I.E. **MINIMIZES SOME LOSS FUNCTION**]
2. REPEAT STEP 1 WITH ANOTHER PARAMETER
3. AT THE END THERE IS A TREE WHERE, AT EACH POINT, ONE OF TWO DECISIONS CAN BE MADE

DECISION TREES (CARTS)

TYPICAL METRICS USED:

[THE IDEA IS TO FIND THE SPLITTING VALUE THAT PUTS ALL OBJECTS OF A GIVEN CLASS IN ONE LEAF]

DECISION TREES (CARTS)

TYPICAL METRIC USED:

GINI IMPURITY

$$G = 1 - \sum_{j=1}^c p_j^2$$

DECISION TREES (CARTS)

TYPICAL METRIC USED:

GINI IMPURITY

$$G = 1 - \sum_{j=1}^c p_j^2$$



fraction of
objects in each class given
a split value

Decision trees (CARTS)

TYPICAL METRIC USED:

GINI IMPURITY

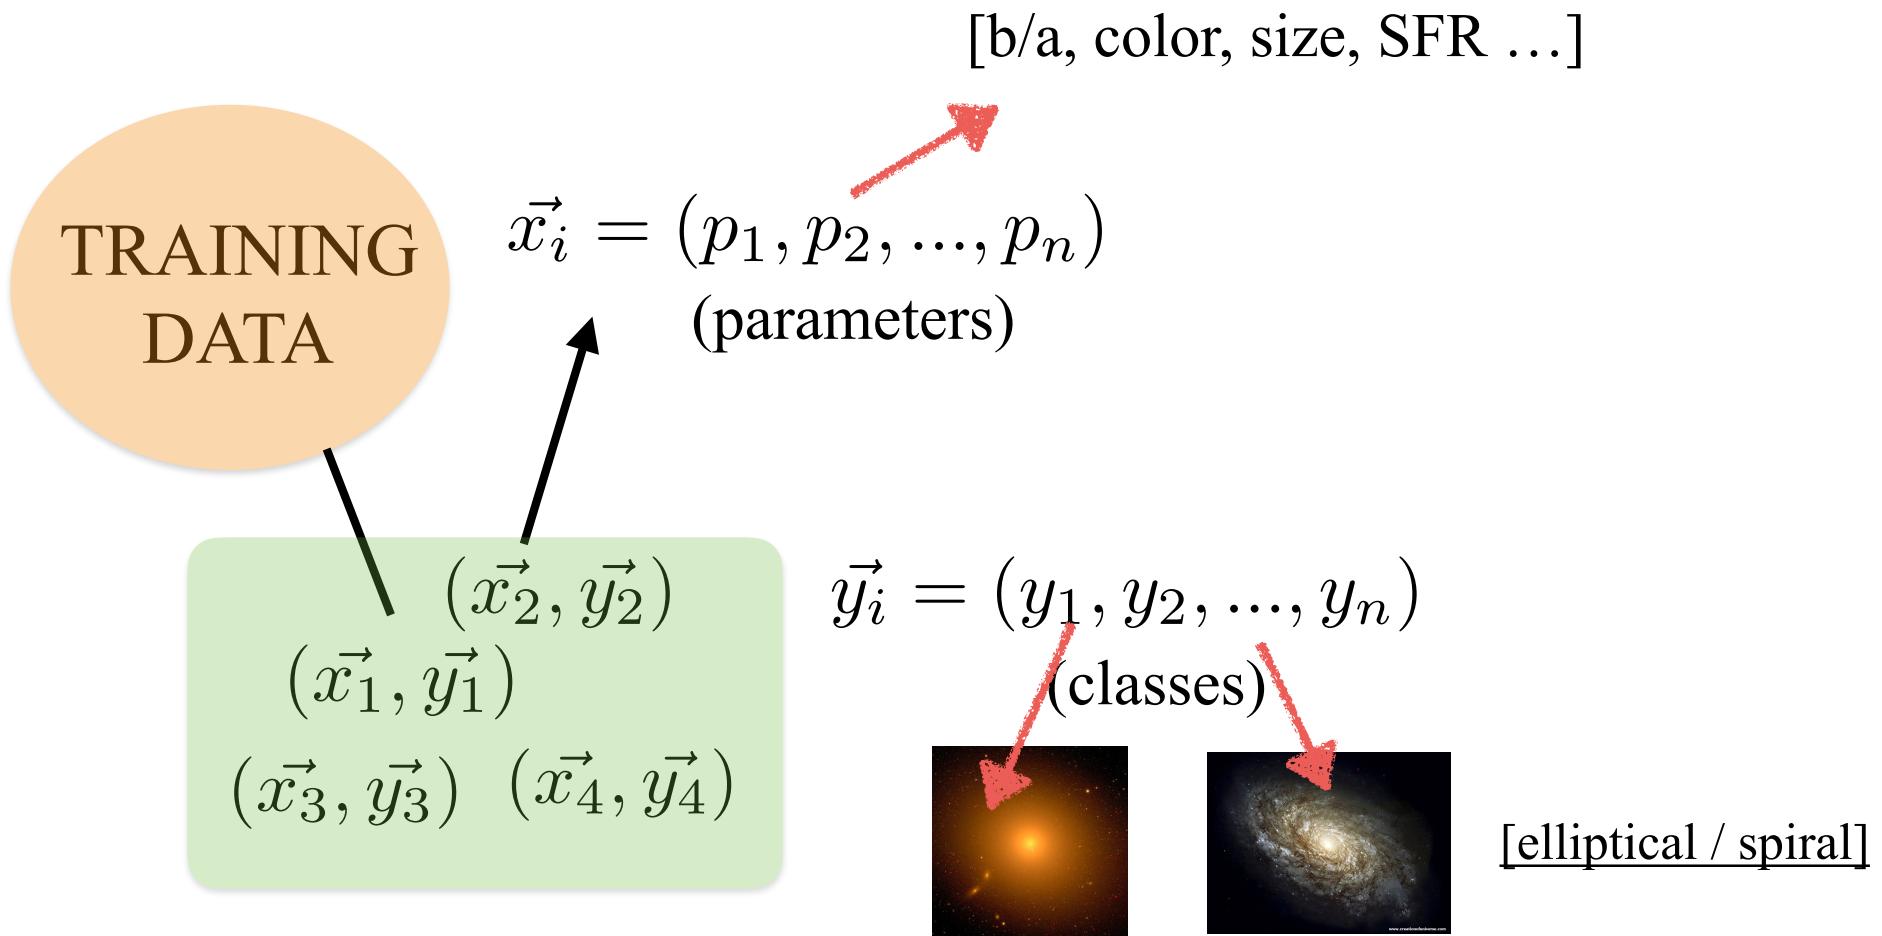
IF THERE IS ONLY ONE CLASS
THE GINI IMPURITY GOES TO 0

$$G = 1 - \sum_{j=1}^c p_j^2$$

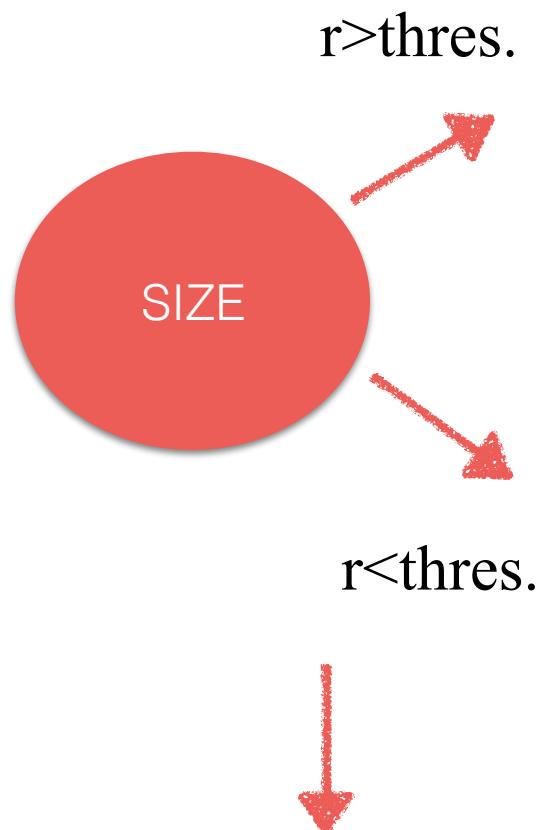


fraction of
objects in each class given
a split value

DECISION TREES (CARTS)



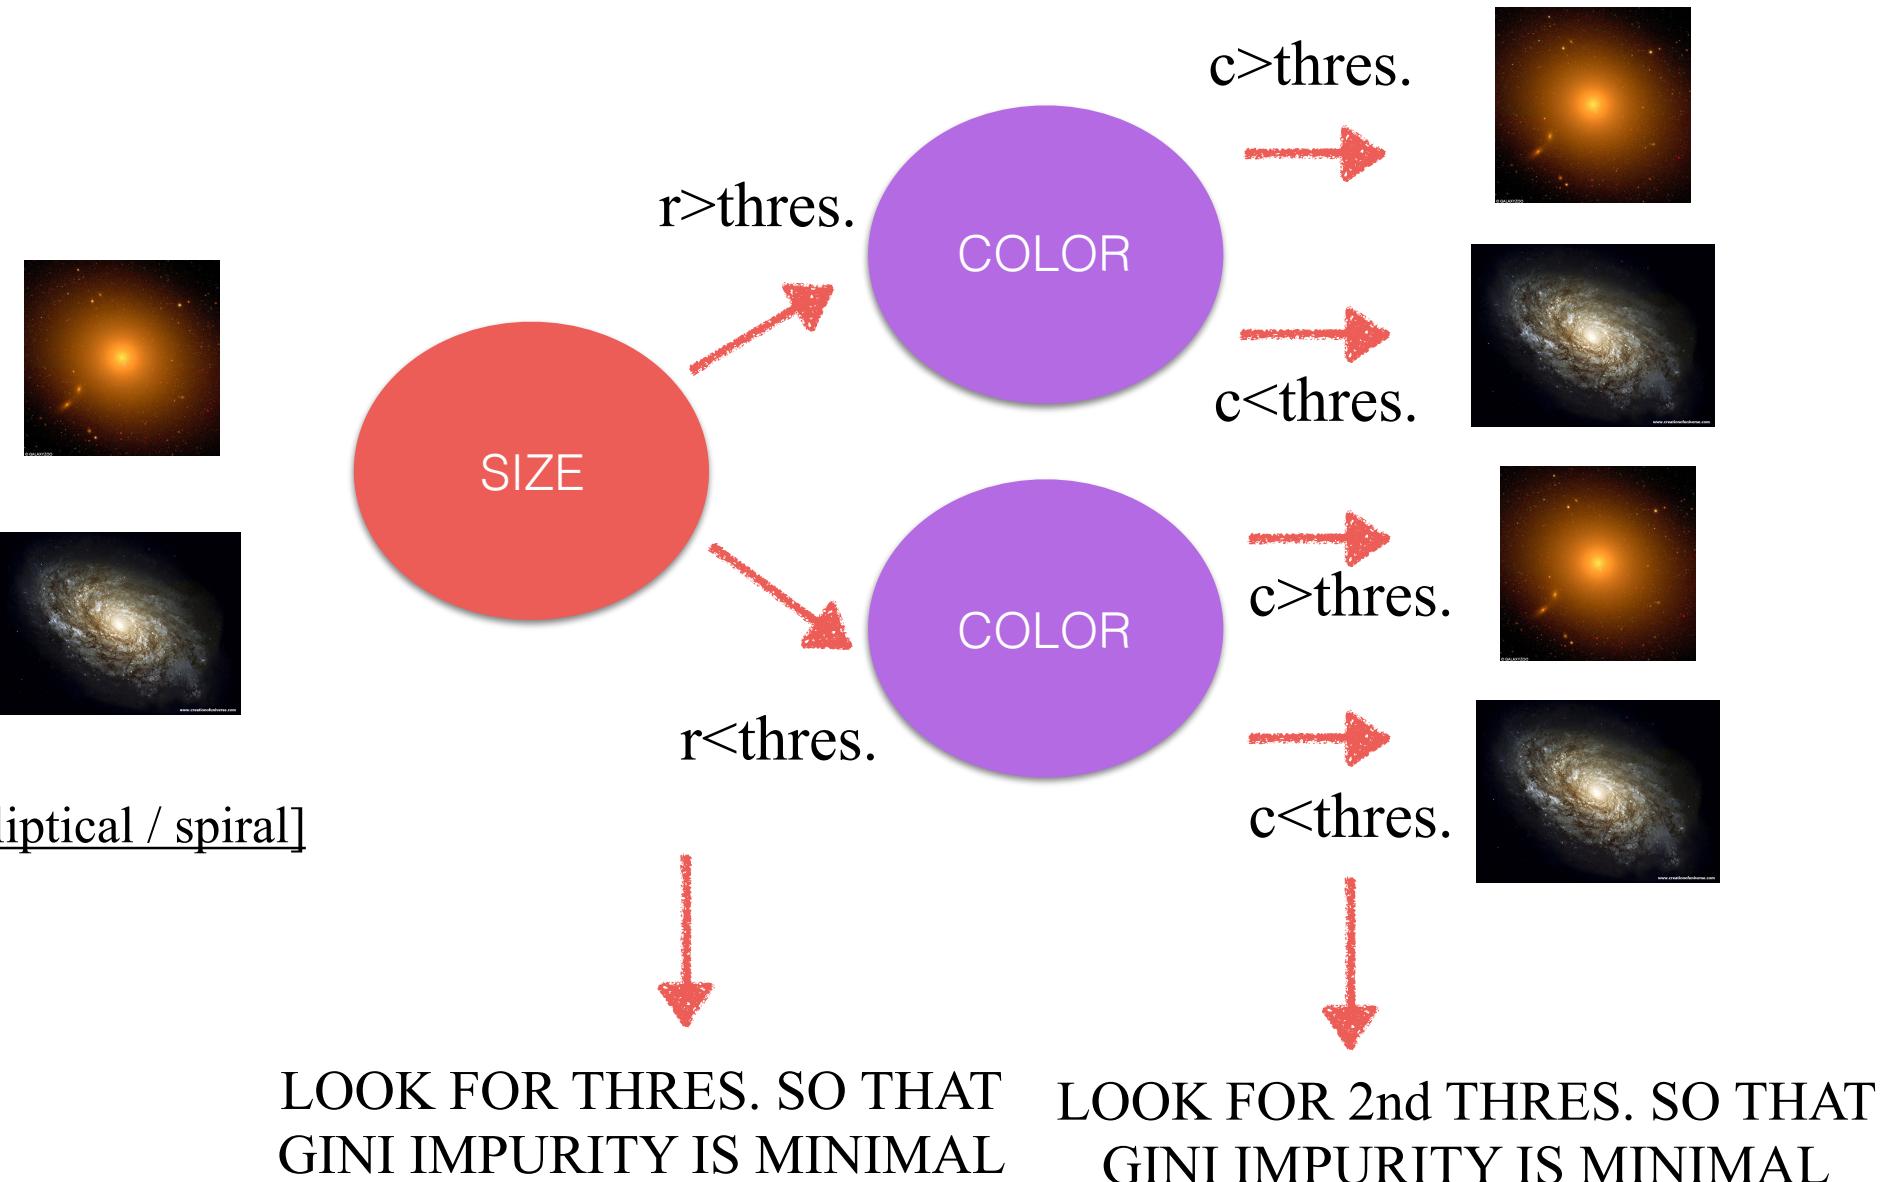
DECISION TREES (CARTS)



[elliptical / spiral]

LOOK FOR THRES. SO THAT
GINI IMPURITY IS MINIMAL

DECISION TREES (CARTS)



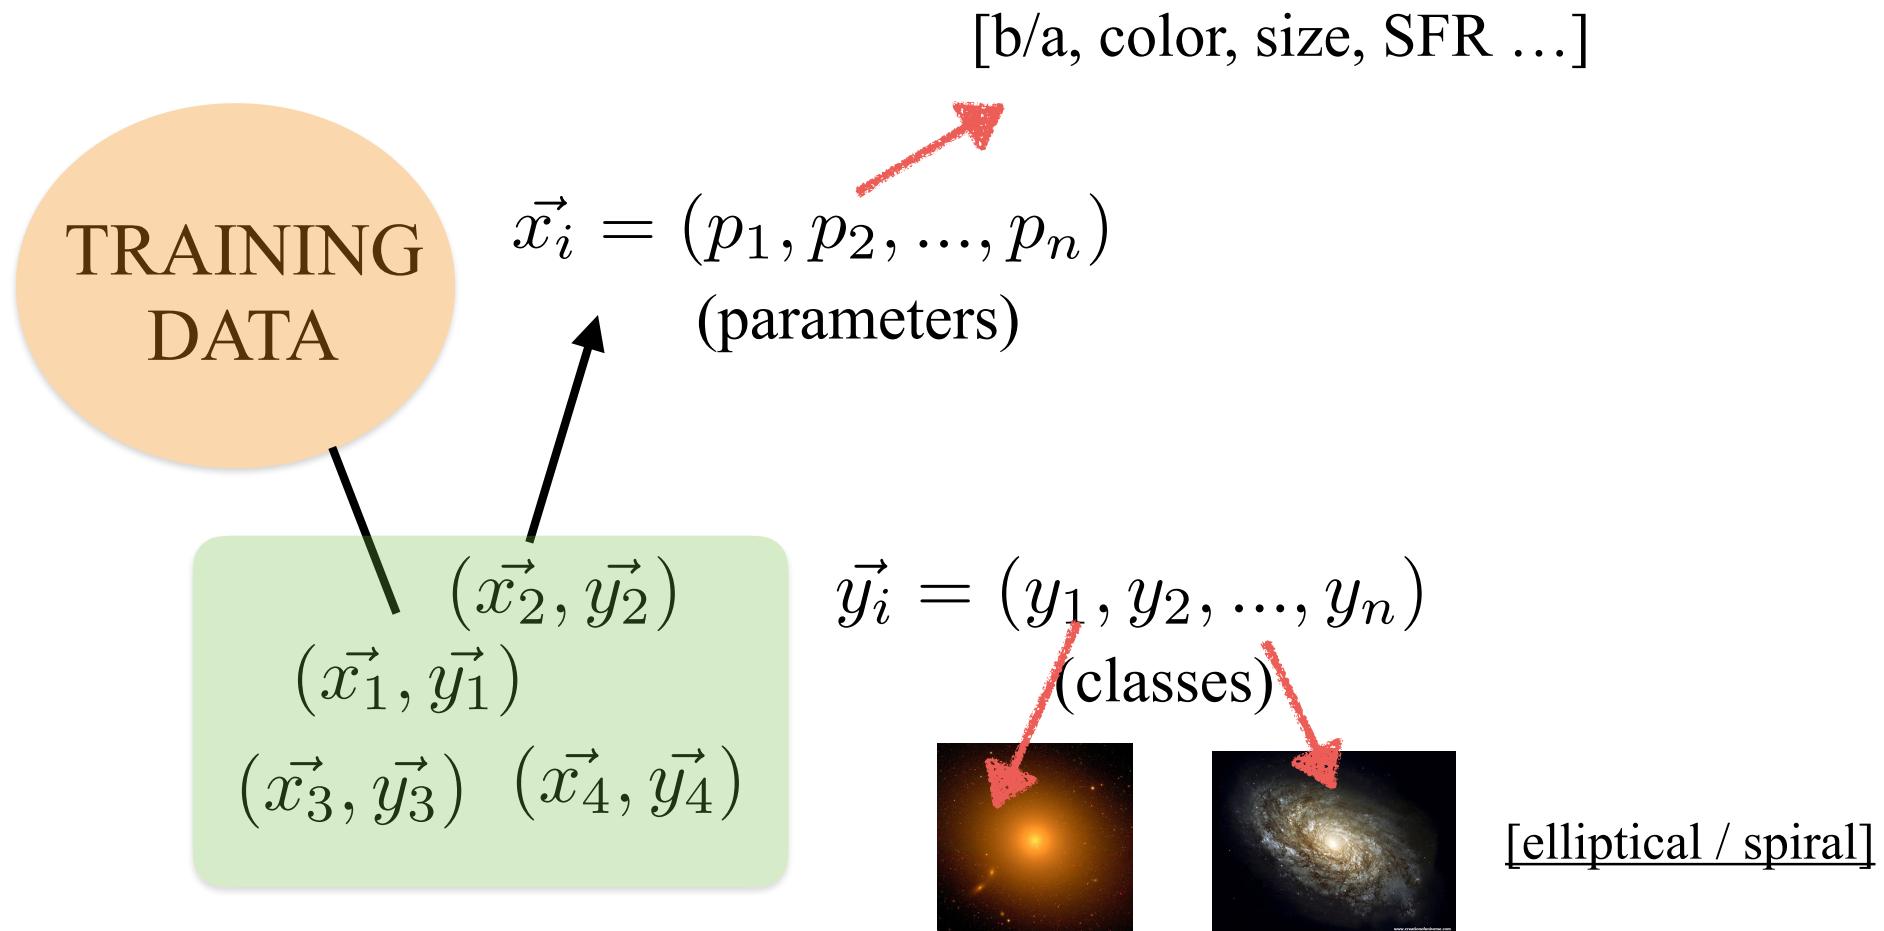
RANDOM FORESTS

ONE PROBLEM WITH CLASSIFICATION TREES IS THAT
THEY CAN EASILY OVERFIT

THE DECISIONS ARE VERY SPECIFIC TO THE TRAINING
SET AND NOT REPRESENTATIVE OF THE FULL
POPULATION

RANDOM FORESTS

RANDOM FORESTS TRY TO SOLVE THIS PROBLEM BY INTRODUCING SOME RANDOM INFORMATION IN THE TRAINING PROCESS



RANDOM FORESTS

(\vec{x}_2, \vec{y}_2)

(\vec{x}_1, \vec{y}_1)

(\vec{x}_3, \vec{y}_3) (\vec{x}_4, \vec{y}_4)

$\vec{x}_i = (p_1, p_2, \dots, p_n)$

RANDOM FORESTS

(\vec{x}_2, \vec{y}_2)
 (\vec{x}_1, \vec{y}_1)
 (\vec{x}_3, \vec{y}_3) (\vec{x}_4, \vec{y}_4)

$\vec{x}_i = (p_1, p_2, \dots, p_n)$



[elliptical / spiral]

SIZE

$r > \text{thres.}$

$r < \text{thres.}$



COLOR

$c > \text{thres.}$



$c < \text{thres.}$



$c > \text{thres.}$



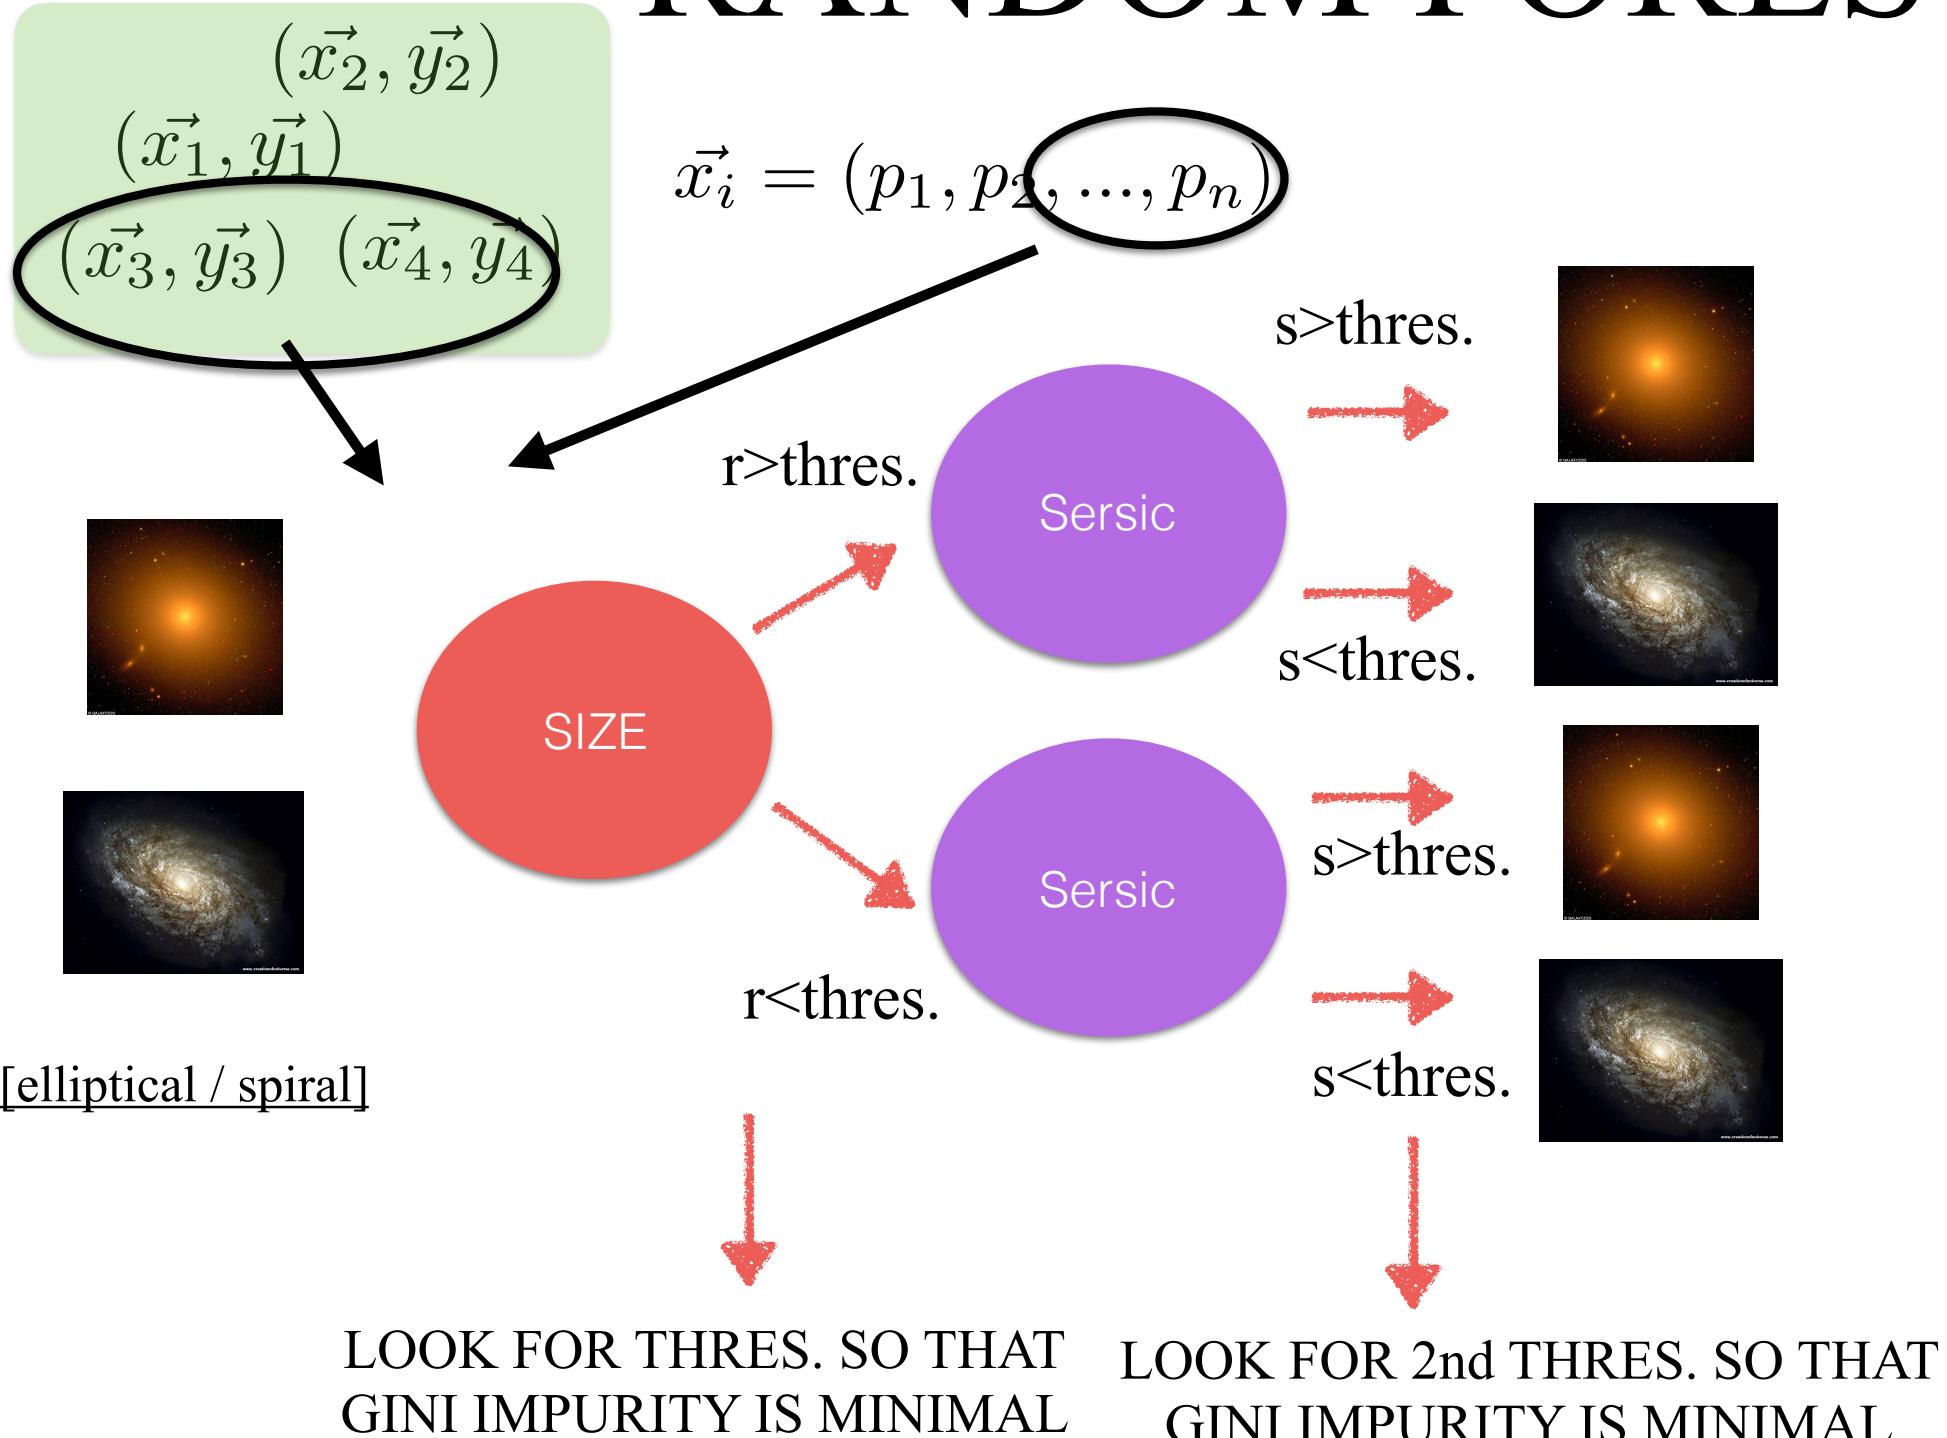
$c < \text{thres.}$



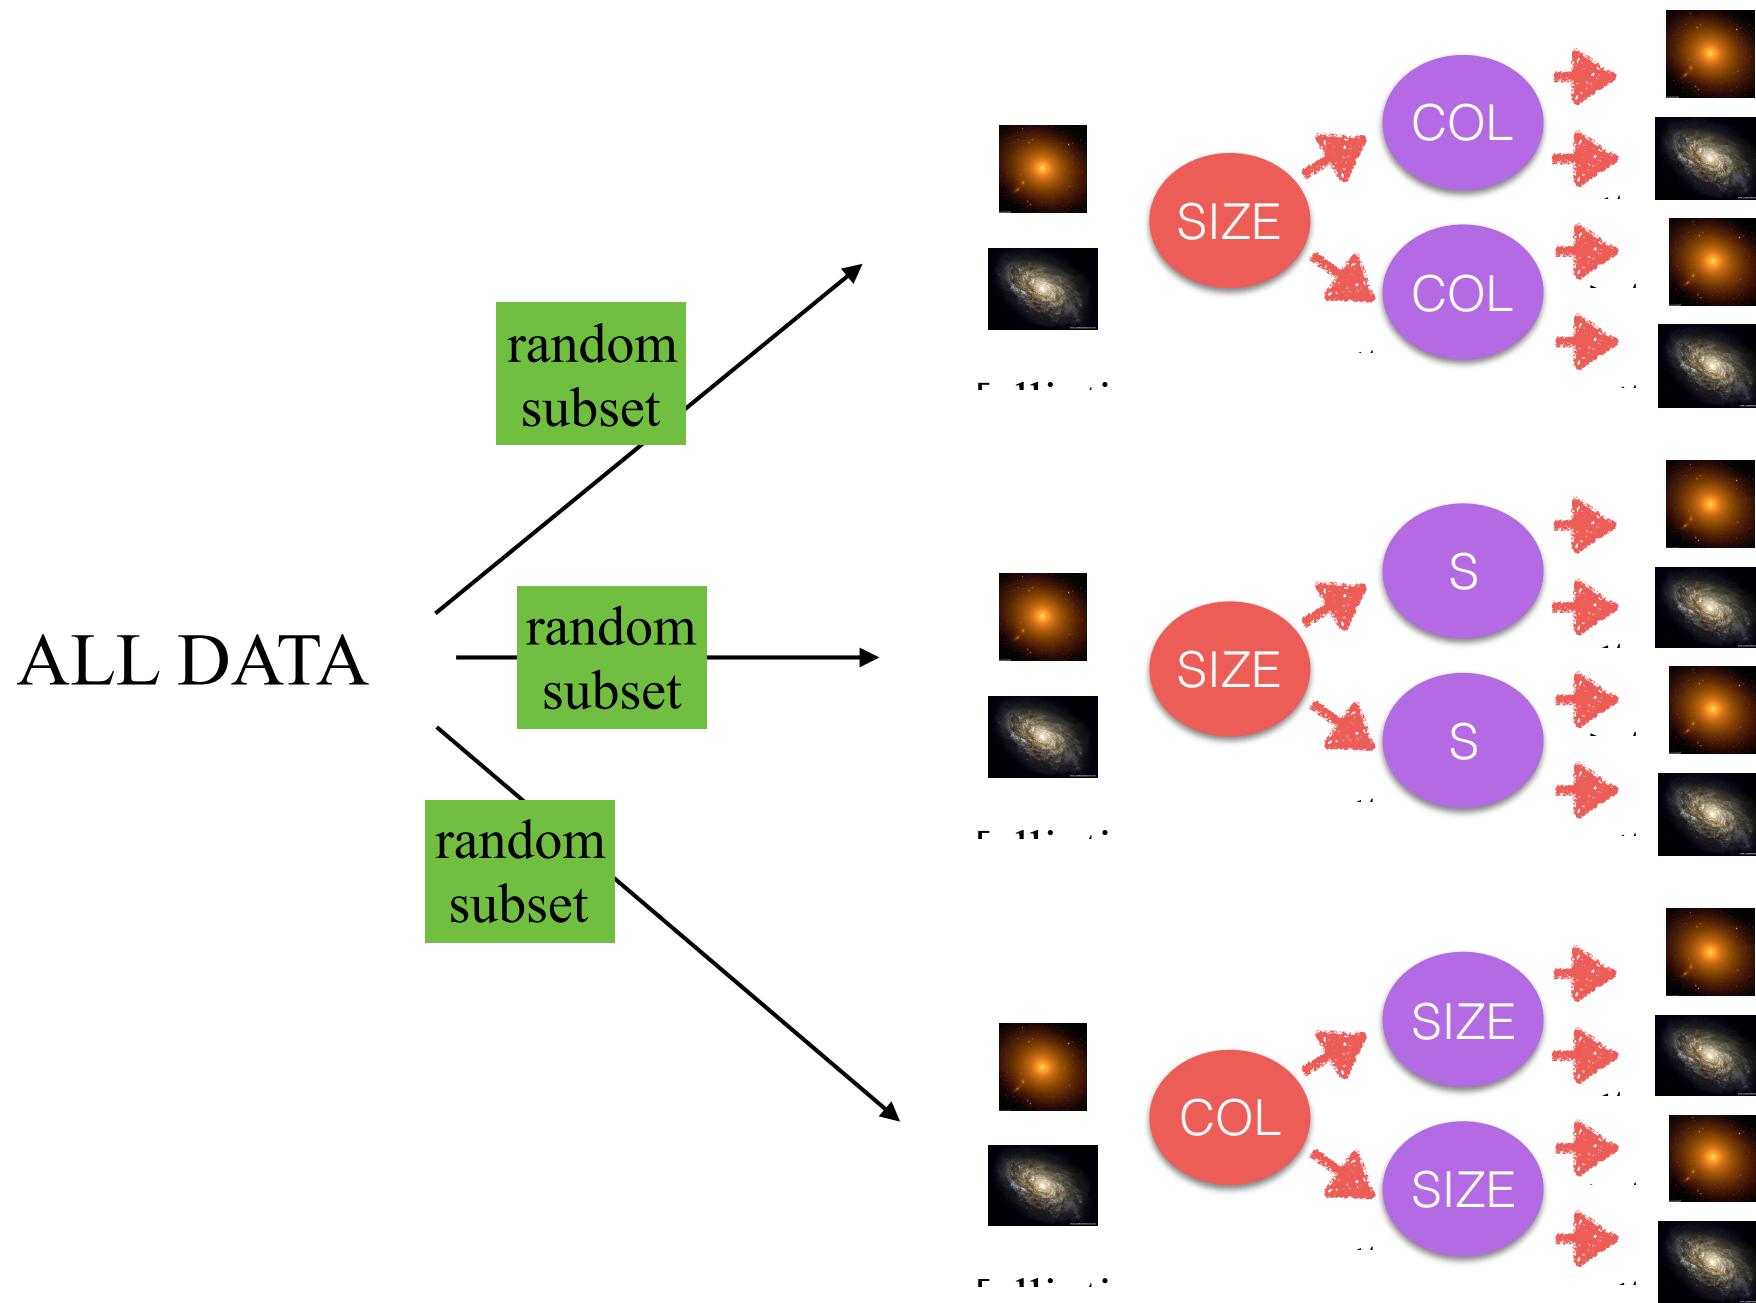
LOOK FOR THRES. SO THAT
GINI IMPURITY IS MINIMAL

LOOK FOR 2nd THRES. SO THAT
GINI IMPURITY IS MINIMAL

RANDOM FORESTS



RANDOM FORESTS



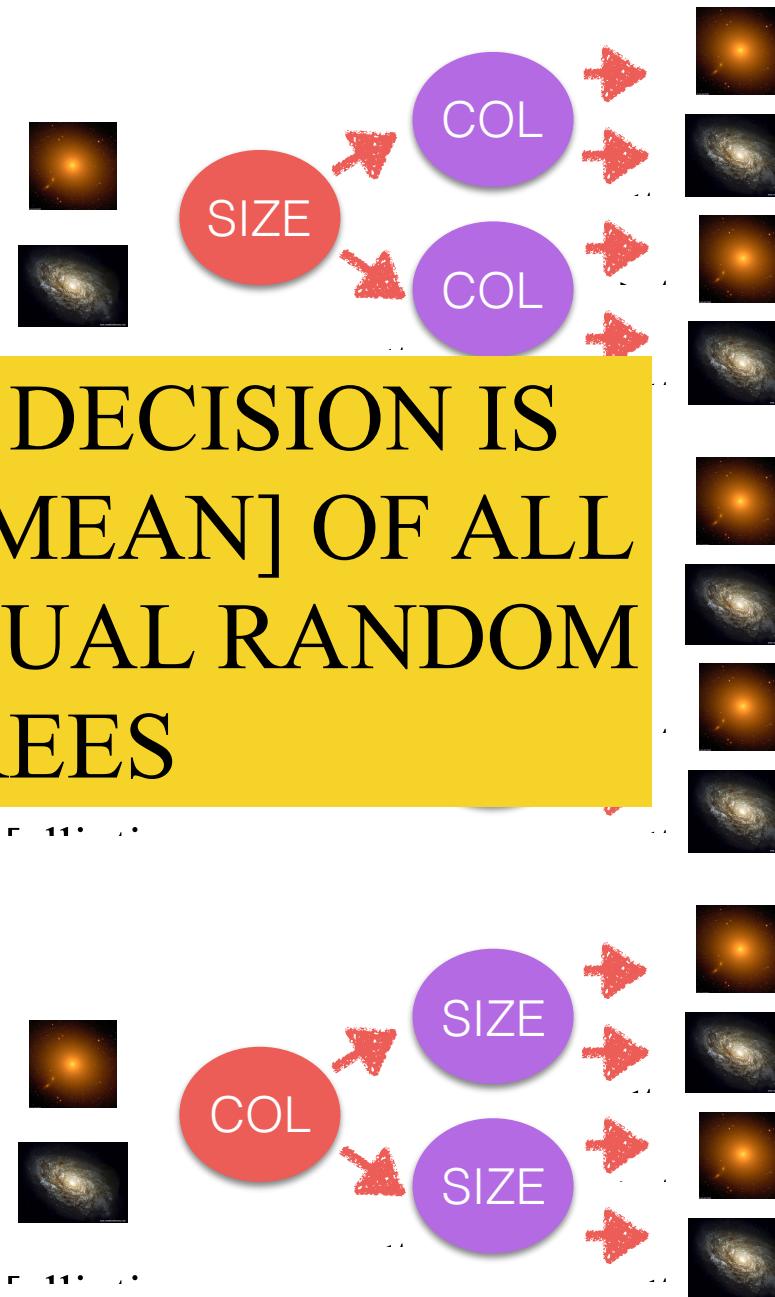
RANDOM FORESTS

ALL DATA

THE FINAL DECISION IS
THE MODE [MEAN] OF ALL
THE INDIVIDUAL RANDOM
TREES

random subset

random



Random Forests

ONE KEY ADVANTAGE OF DECISION TREE ALGORITHMS IS THAT
THEY ARE VERY EASY TO INTERPRET

ONE CAN EASILY DETERMINE THE MOST IMPORTANT FEATURES
TO TAKE DECISIONS

RANDOM FORESTS

ONE KEY ADVANTAGE OF DECISION TREE ALGORITHMS IS THAT THEY ARE VERY EASY TO INTERPRET

ONE CAN EASILY DETERMINE THE MOST IMPORTANT FEATURES TO TAKE DECISIONS

Rank	Property	AUC	Success label ^a
1	ALL	0.9074 ± 0.0106	Outstanding
1	CVD	0.8559 ± 0.0039	Excellent
2	M_{bulge}	0.8335 ± 0.0060	Excellent
3	B/T	0.8267 ± 0.0028	Excellent
4	M_{halo}	0.7983 ± 0.0045	Acceptable
5	M_*	0.7819 ± 0.0025	Acceptable
6	M_{disc}	0.7124 ± 0.0016	Acceptable
7	δ_5	0.5894 ± 0.0015	Unacceptable
8	Re	0.5599 ± 0.0013	Unacceptable

IMPORTANCE OF
PARAMETERS TO PREDICT
IF A GALAXY IS QUENCHED

PRACTICAL NOTE: *Python scikit learn*

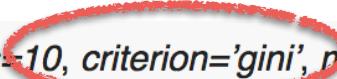
- All these different methods are standard and available in Python
- Very easy to use
- All info here

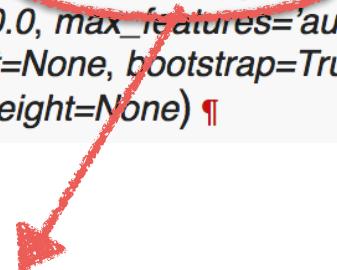
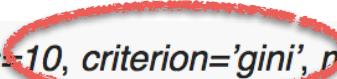
sklearn.ensemble.RandomForestClassifier

```
class sklearn.ensemble. RandomForestClassifier(n_estimators=10, criterion='gini', max_depth=None,  
min_samples_split=2, min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features='auto',  
max_leaf_nodes=None, min_impurity_decrease=0.0, min_impurity_split=None, bootstrap=True, oob_score=False,  
n_jobs=1, random_state=None, verbose=0, warm_start=False, class_weight=None)  [source]
```

NUMBER OF RANDOM
TREES

sklearn.ensemble.RandomForestClassifier

```
class sklearn.ensemble. RandomForestClassifier (n_estimators=10, criterion='gini', max_depth=None,  
min_samples_split=2, min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features='auto',  
max_leaf_nodes=None, min_impurity_decrease=0.0, min_impurity_split=None, bootstrap=True, oob_score=False,  
n_jobs=1, random_state=None, verbose=0, warm_start=False, class_weight=None)  [source]
```



METRIC TO MINIMIZE

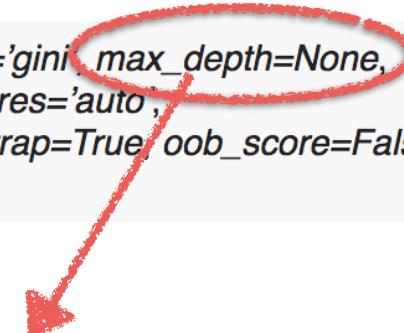
sklearn.ensemble.RandomForestClassifier

```
class sklearn.ensemble. RandomForestClassifier (n_estimators=10, criterion='gini', max_depth=None,  
min_samples_split=2, min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features='auto',  
max_leaf_nodes=None, min_impurity_decrease=0.0, min_impurity_split=None, bootstrap=True, oob_score=False,  
n_jobs=1, random_state=None, verbose=0, warm_start=False, class_weight=None) ¶ [source]
```

MAXIMUM NUMBER OF SPLITS

sklearn.ensemble.RandomForestClassifier

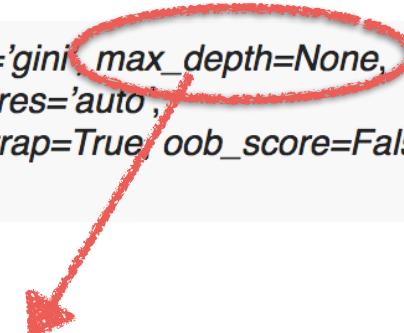
```
class sklearn.ensemble.RandomForestClassifier(n_estimators=10, criterion='gini', max_depth=None,  
min_samples_split=2, min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features='auto',  
max_leaf_nodes=None, min_impurity_decrease=0.0, min_impurity_split=None, bootstrap=True, oob_score=False,  
n_jobs=1, random_state=None, verbose=0, warm_start=False, class_weight=None) ¶ [source]
```



MAXIMUM NUMBER OF SPLITS

```
>>> from sklearn.ensemble import RandomForestClassifier  
>>> from sklearn.datasets import make_classification  
>>>  
>>> X, y = make_classification(n_samples=1000, n_features=4,  
...                                n_informative=2, n_redundant=0,  
...                                random_state=0, shuffle=False)  
>>> clf = RandomForestClassifier(max_depth=2, random_state=0)  
>>> clf.fit(X, y)  
RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',  
                      max_depth=2, max_features='auto', max_leaf_nodes=None,  
                      min_impurity_decrease=0.0, min_impurity_split=None,  
                      min_samples_leaf=1, min_samples_split=2,  
                      min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=1,  
                      oob_score=False, random_state=0, verbose=0, warm_start=False)  
>>> print(clf.feature_importances_)  
[ 0.17287856  0.80608704  0.01884792  0.00218648]  
>>> print(clf.predict([[0, 0, 0, 0]]))  
[1]
```

sklearn.ensemble.RandomForestClassifier

```
class sklearn.ensemble.RandomForestClassifier(n_estimators=10, criterion='gini', max_depth=None,  
min_samples_split=2, min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features='auto',  
max_leaf_nodes=None, min_impurity_decrease=0.0, min_impurity_split=None, bootstrap=True, oob_score=False,  
n_jobs=1, random_state=None, verbose=0, warm_start=False, class_weight=None)  [source]
```

MAXIMUM NUMBER OF SPLITS

```
>>> from sklearn.ensemble import RandomForestClassifier  
>>> from sklearn.datasets import make_classification  
>>>  
>>> X, y = make_classification(n_samples=1000, n_features=4,  
...                                n_informative=2, n_redundant=0,  
...                                random_state=0, shuffle=False)  
>>> clf = RandomForestClassifier(max_depth=2, random_state=0)  
>>> clf.fit(X, y)  
RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',  
                      max_depth=2, max_features='auto', max_leaf_nodes=None,  
                      min_impurity_decrease=0.0, min_impurity_split=None,
```

DIFFERENT CLASSIFIERS ARE OBJECTS WITH METHODS TO
FIT, PREDICT ETC..

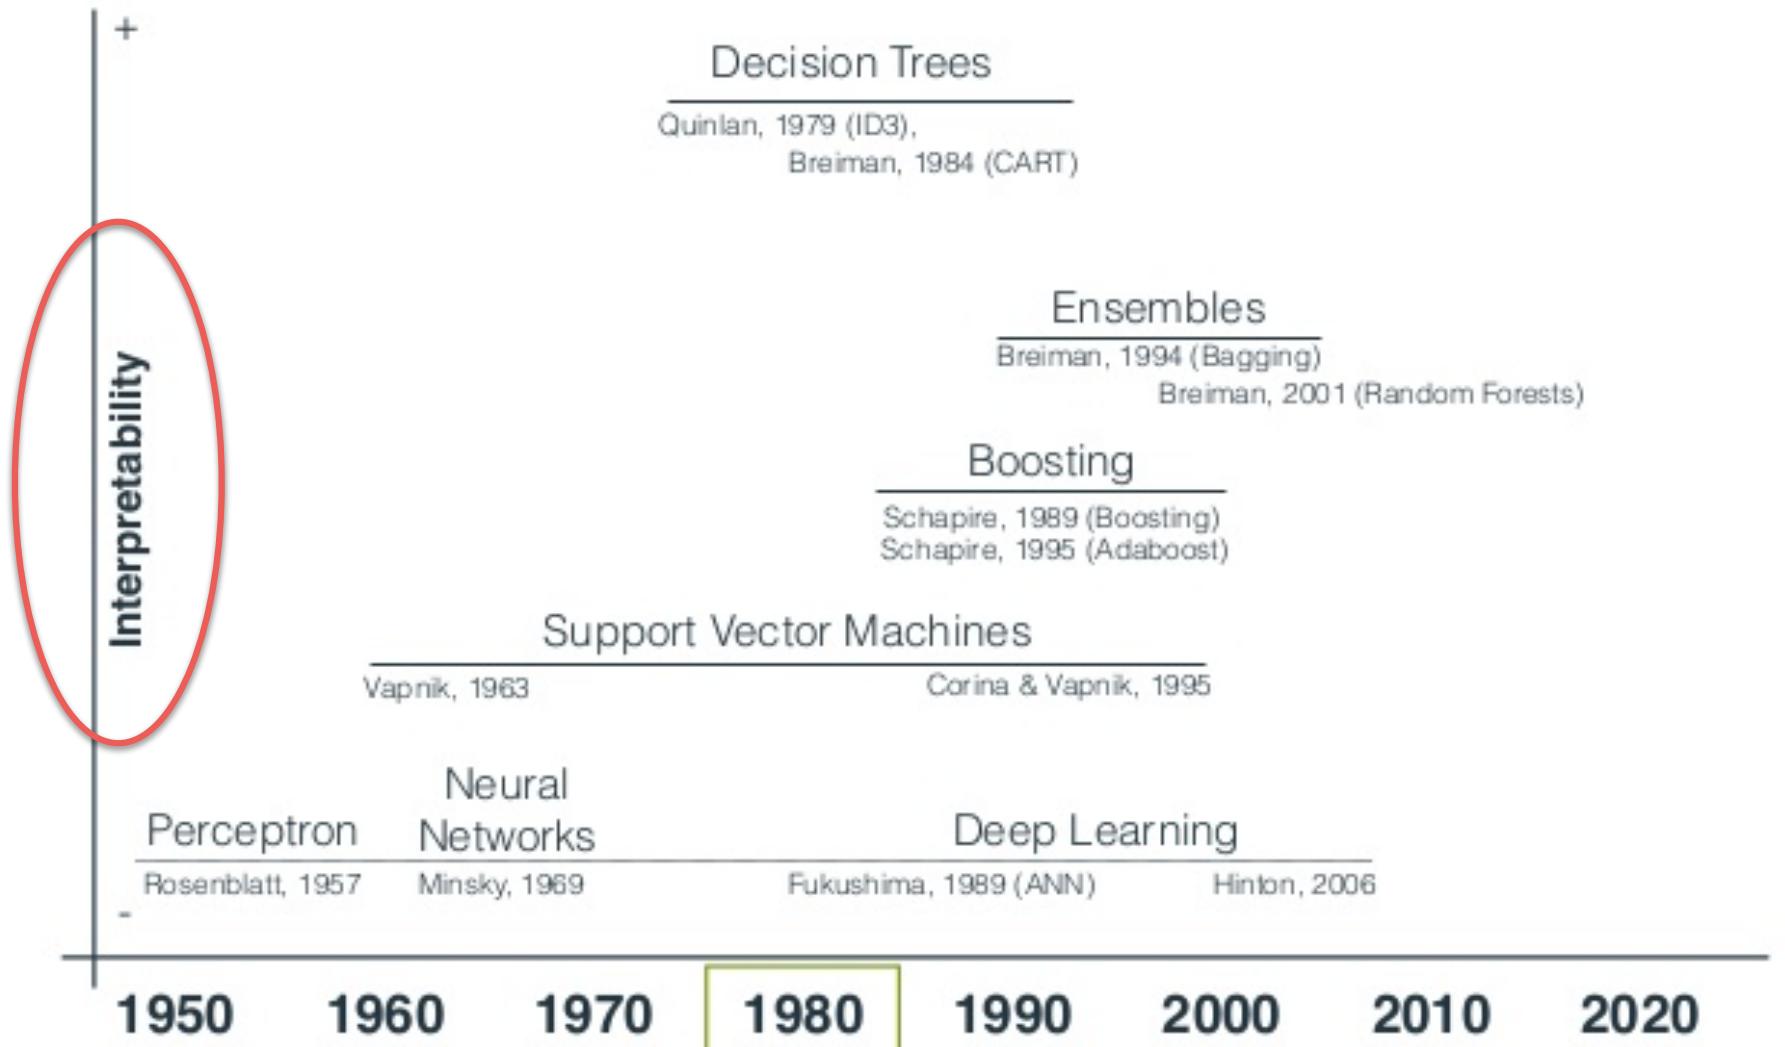
```
>>> print(clf.predict([[0, 0, 0, 0]]))  
[1]
```

HOW TO CHOOSE YOUR CLASSICAL CLASSIFIER?

NO RULE OF THUMB - REALLY DEPENDS ON APPLICATION

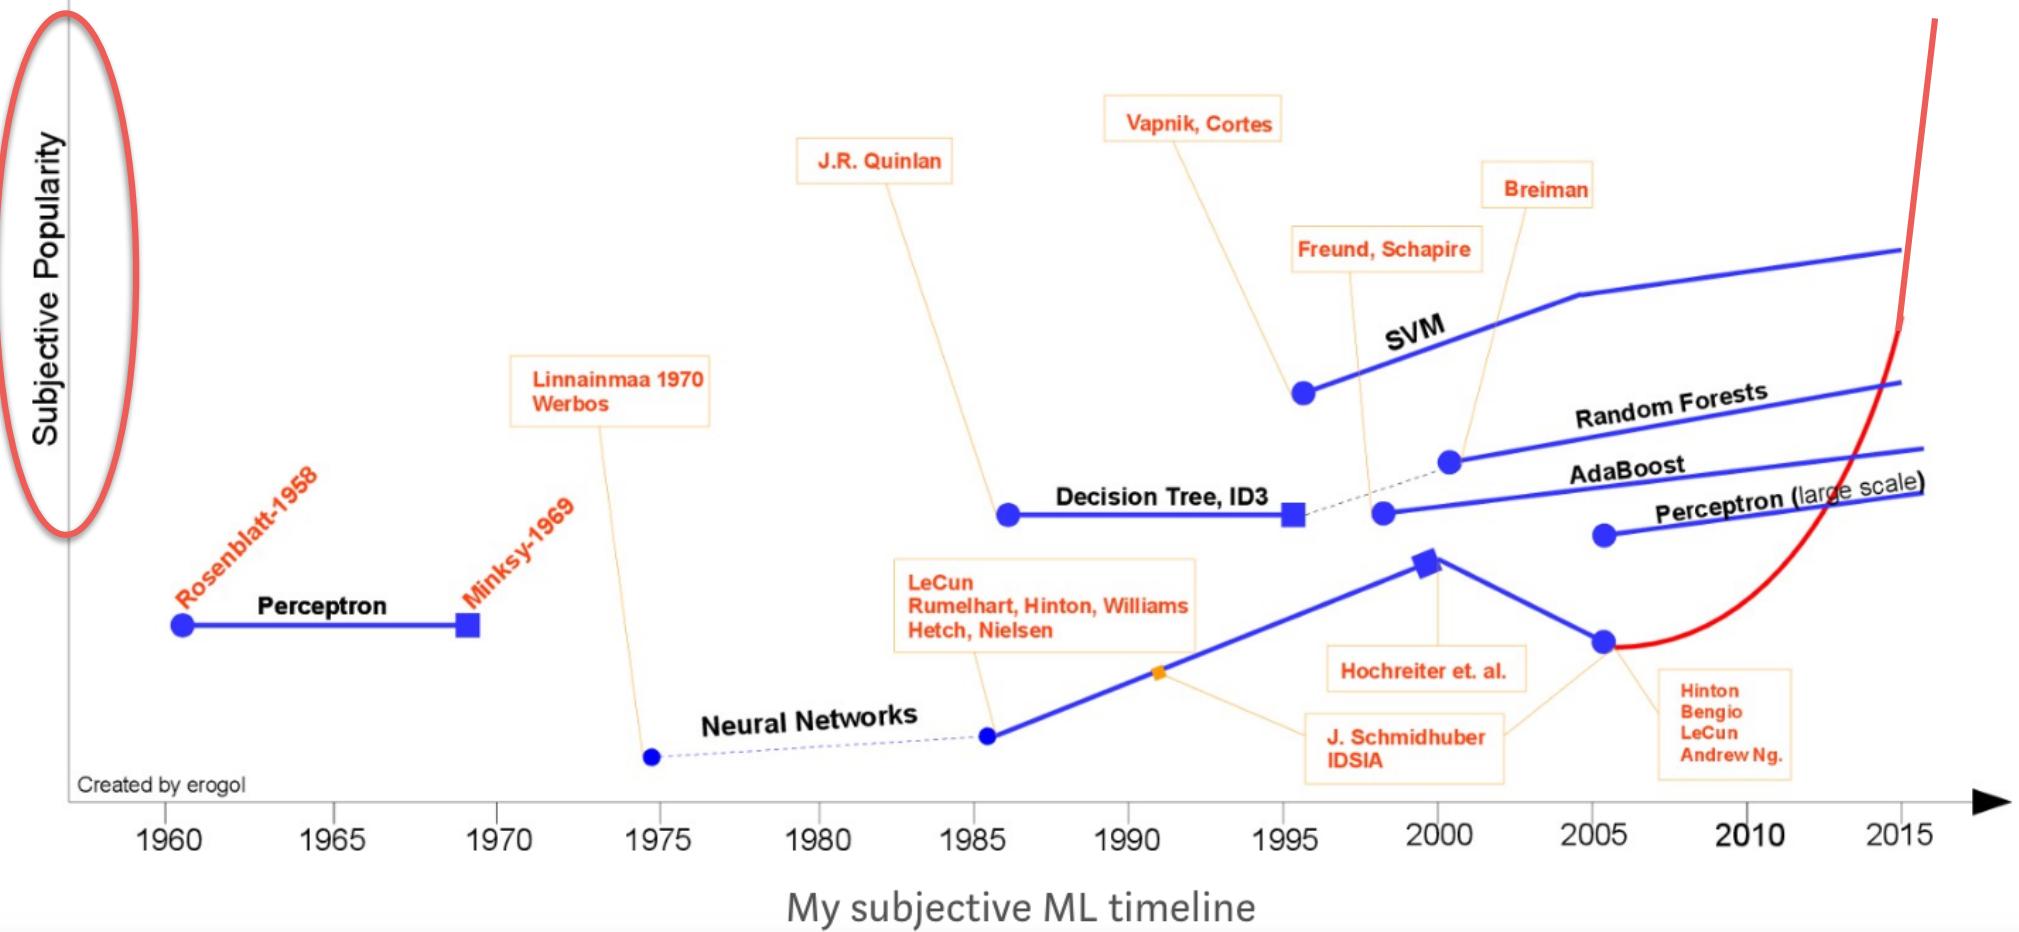
ML METHOD	++	-	Python
CARTS / RANDOM FOREST	Easy to interpret (“White box”) Litte data preparation Both numerical + categorical	Over-complex trees Unstable Biased trees if some classes dominate	sklearn.ensemble.RandomForestClassifier sklearn.ensemble.RandomForestRegressor
SVM	Easy to interpret + Fast Kernel trick allows no linear problems	not very well suited to multi-class problems	sklearn.svm sklearn.svc
NN	seed of deep-learning very efficient with large amount of data as we will see	more difficult to interpret computing intensive	sklearn.neural_network.MLPClassifier sklearn.neural_network.MLPRegressor

CAN DEPEND ON YOUR MAIN INTEREST



credit

ALSO INFLUENCED BY “MAINSTREAM” TRENDS

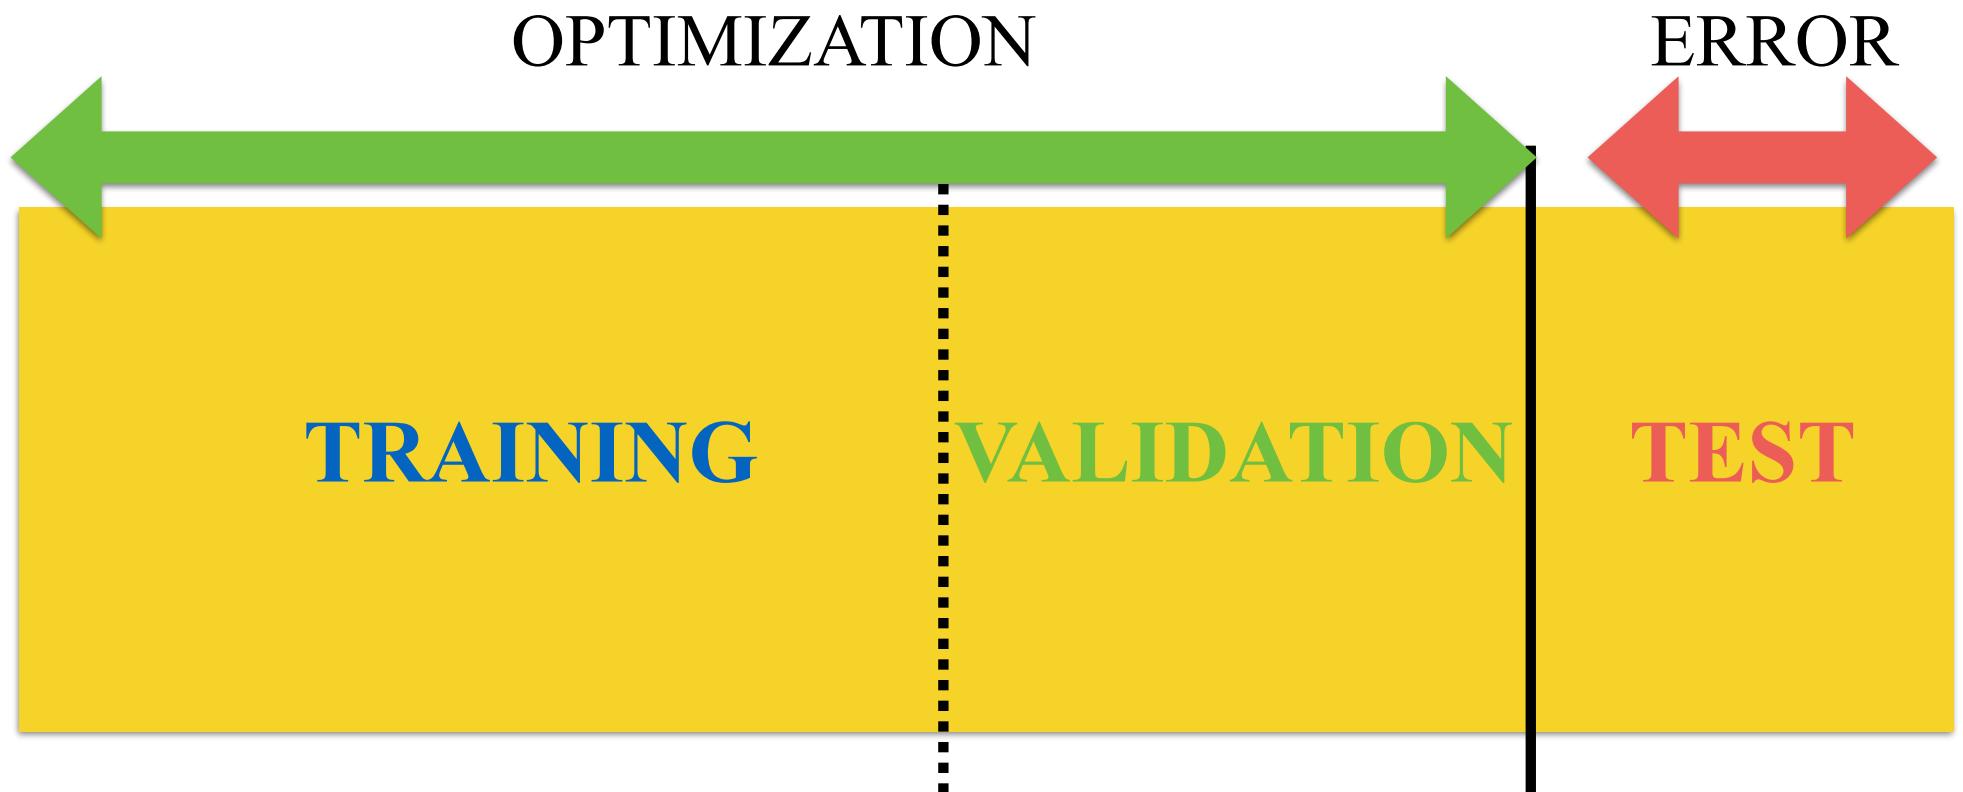


Source

HOW DO I KNOW THAT MY MACHINE LEARNING ALGORITHM IS WORKING?

- The way results are evaluated depends strongly on the type of problem
 - Binary Classification: ACC, ROC, P-C
 - Multi-Class: Confusion Matrix
 - Regression: RMSE, Bias, Scatter etc..

REMEMBER



training set: use to train the classifier

validation set: use to monitor performance in real time - check
for overfitting

test set: use to train the classifier

Evaluation of results [binary class.]

THE MOST STRAIGHTFORWARD WAY IS TO EVALUATE THE
TOTAL ACCURACY

$$ACC = \frac{TP + TN}{TP + TN + FP + FN}$$

Annotations: "true positives" points to the TP term in the numerator; "true negatives" points to the TN term in the numerator.

MEASURES HOW MANY OBJECTS ARE CORRECTLY
CLASSIFIED

Evaluation of results [binary class.]

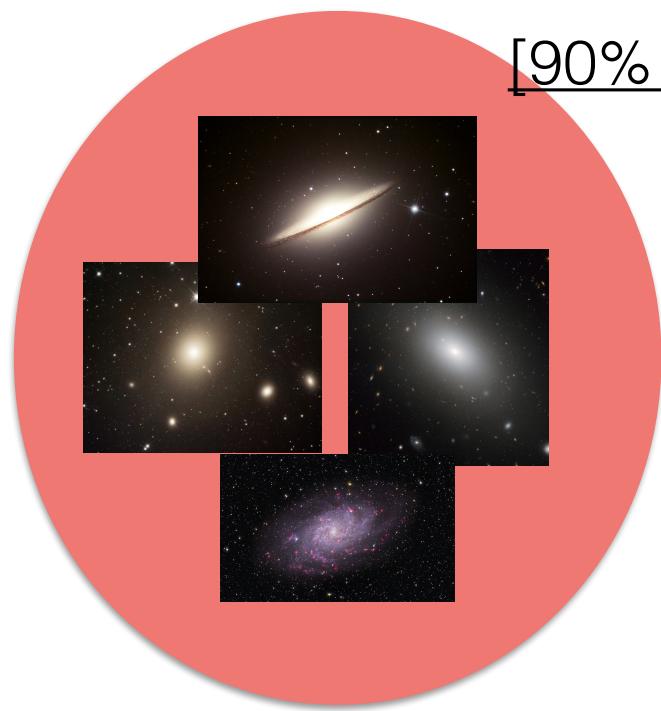
THE MOST STRAIGHTFORWARD WAY IS TO EVALUATE THE
TOTAL ACCURACY

$$ACC = \frac{\text{true positives} + \text{true negatives}}{TP + TN + FP + FN}$$

NOT VERY
INFORMATIVE IF
UNBALANCED
CLASSES

MEASURES HOW MANY OBJECTS ARE CORRECTLY
CLASSIFIED

IMAGINE AN EXTREME CASE WITH VERY UNBALANCED
DATA...



[90% E, 10% S]